

Markov Chain Monte Carlo (MCMC)

Tim Frasier

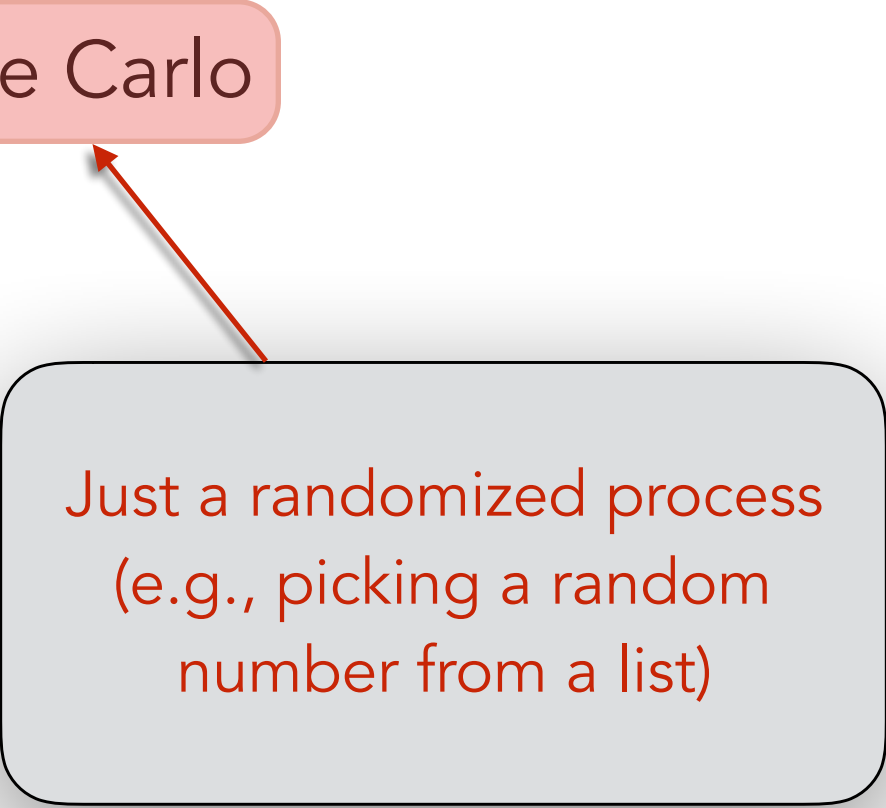
What is MCMC?

What Is MCMC?

- Markov Chain Monte Carlo

What Is MCMC?

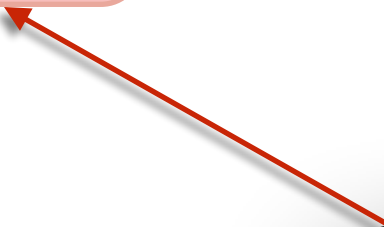
- Markov Chain Monte Carlo



Just a randomized process
(e.g., picking a random
number from a list)

What Is MCMC?

- Markov Chain Monte Carlo



A process where the next step only depends on the current location (this will make more sense in a minute)

What Is MCMC?

Think of a robot walking through a field

- Step 1: Pick a random point along an edge to start at
- Step 2: Randomly select what direction it turns
- Step 3: If step will keep robot in field, take one step in that direction

If steps 2 & 3 repeated enough times, will eventually visit every position in the field even though it is a random process

What Is MCMC?

Think of a robot walking through a field

- Step 1: Pick a random point along an edge to start at
- Step 2: Randomly select what direction it turns
- Step 3: If step will keep robot in field, take one step in that direction

If steps 2 & 3 repeated enough times, will eventually visit every point in the field, even though it is a

The Monte Carlo part

What Is MCMC?

Think of a robot walking through a field

- Step 1: Pick a random point along an edge to start at
- Step 2: Randomly select what direction it turns
- Step 3: If step will keep robot in field, take one step in that direction


If steps 2 & 3 repeated enough times, will eventually visit every position in the field even though it is a

The Markov Chain part
(where it goes next depends
on where it is now)

What Is MCMC?

Think of a robot walking through a field

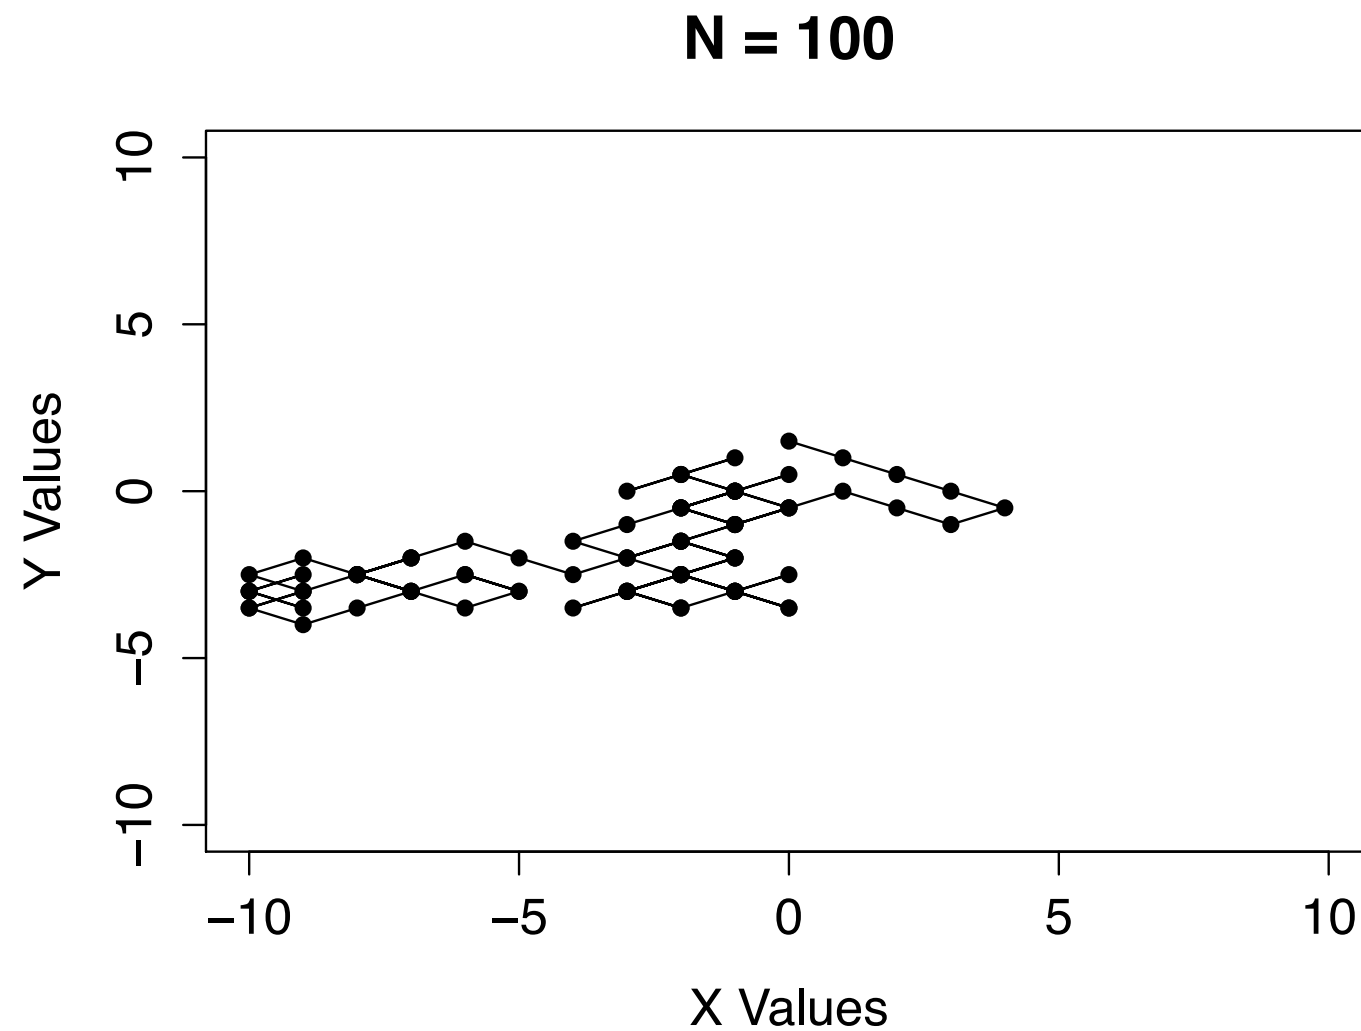
- 100 steps



What Is MCMC?

Think of a robot walking through a field

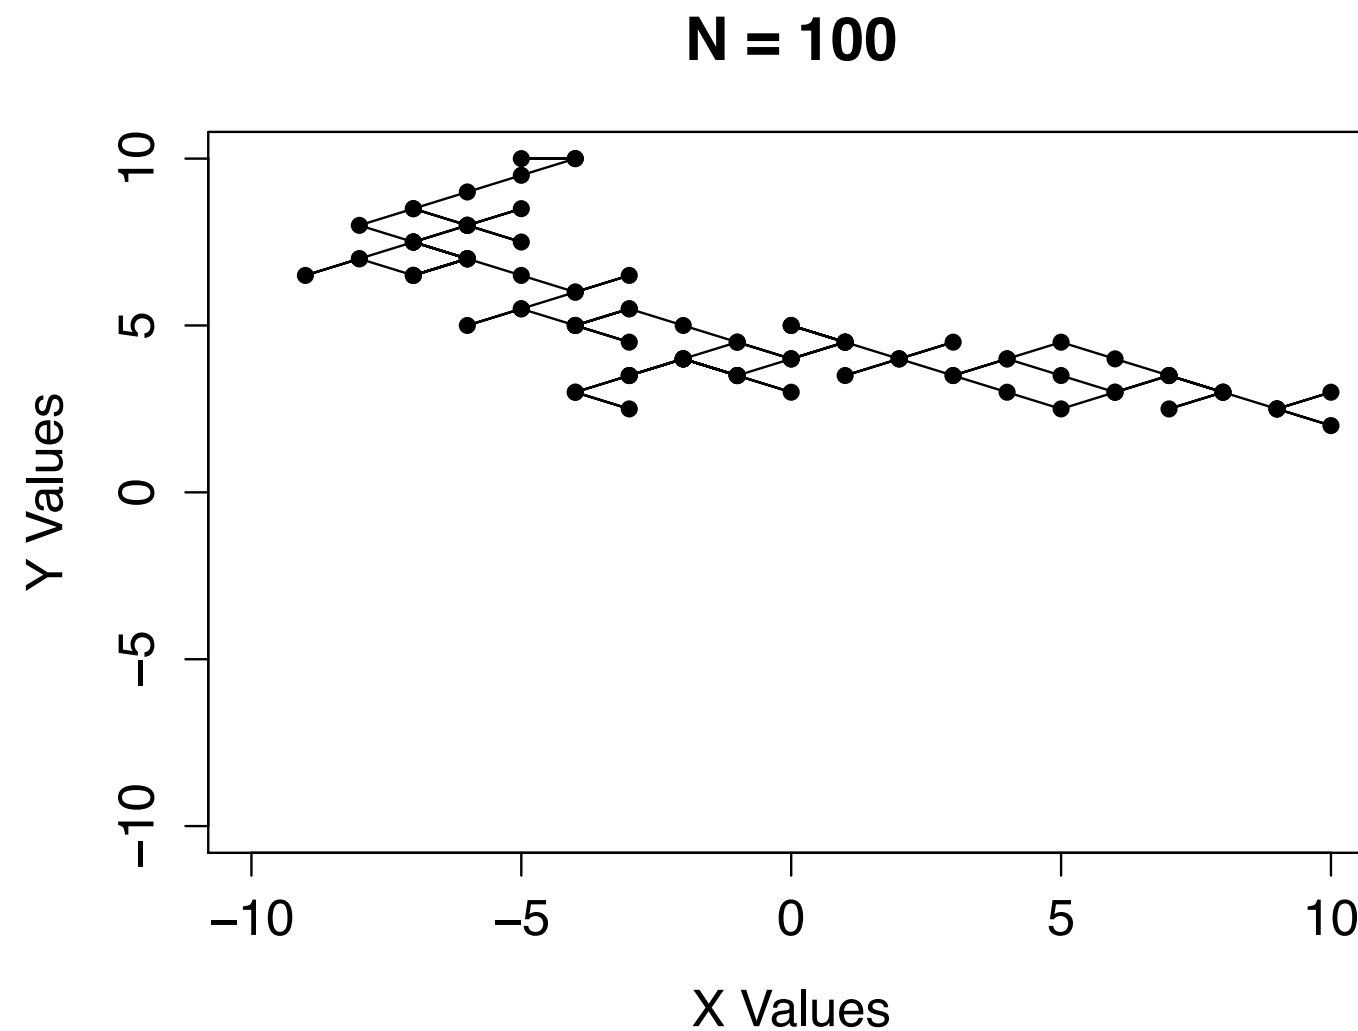
- 100 steps (again)



What Is MCMC?

Think of a robot walking through a field

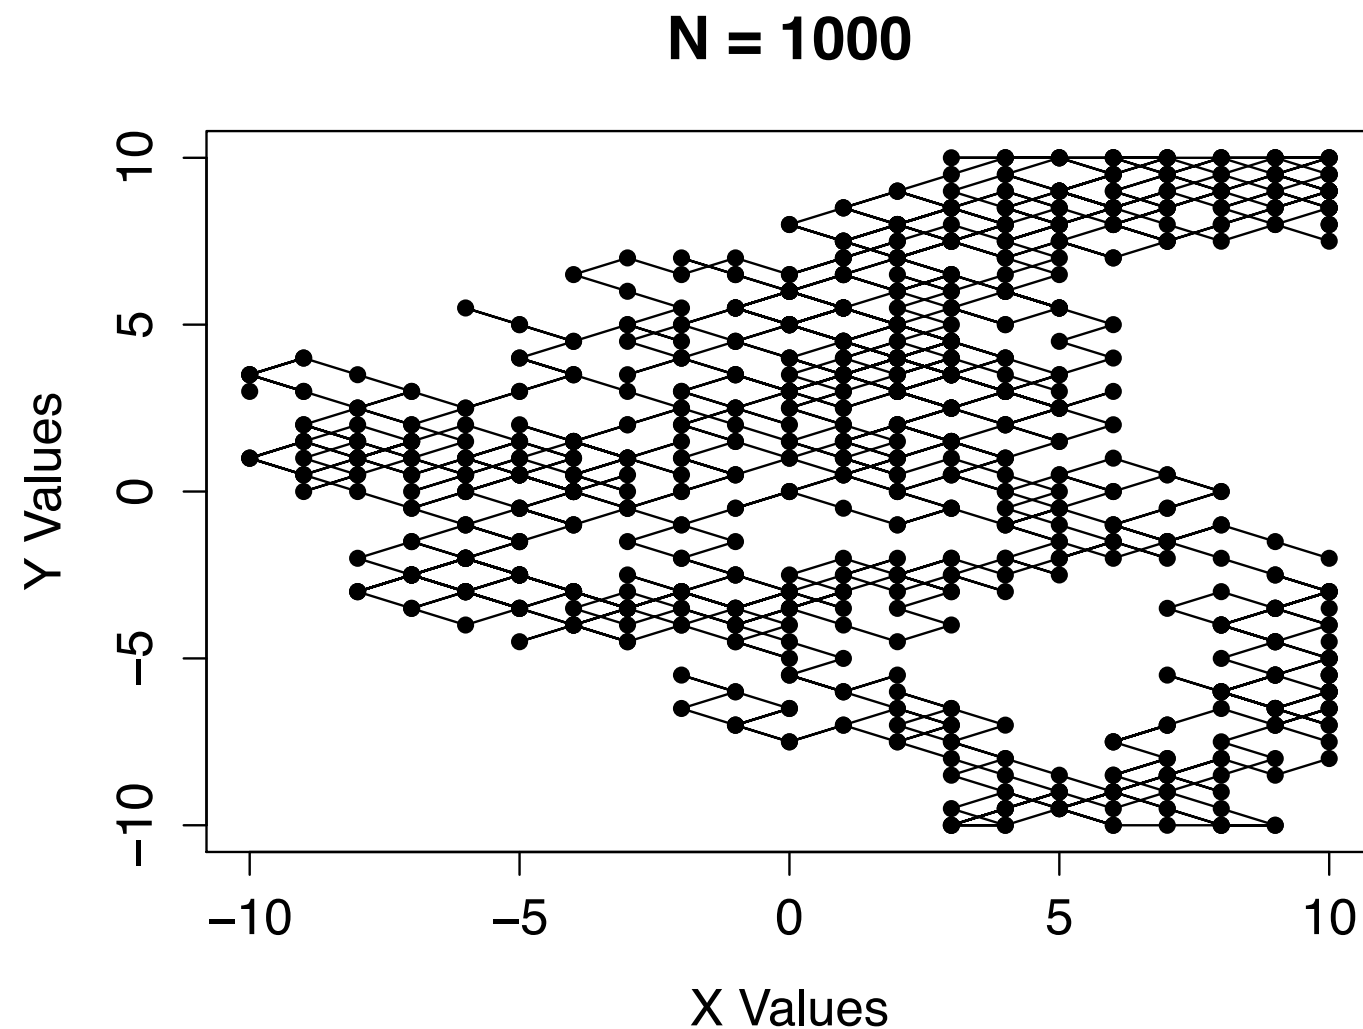
- 100 steps (and again)



What Is MCMC?

Think of a robot walking through a field


- 1,000 steps



What Is MCMC?

Think of a robot walking through a field

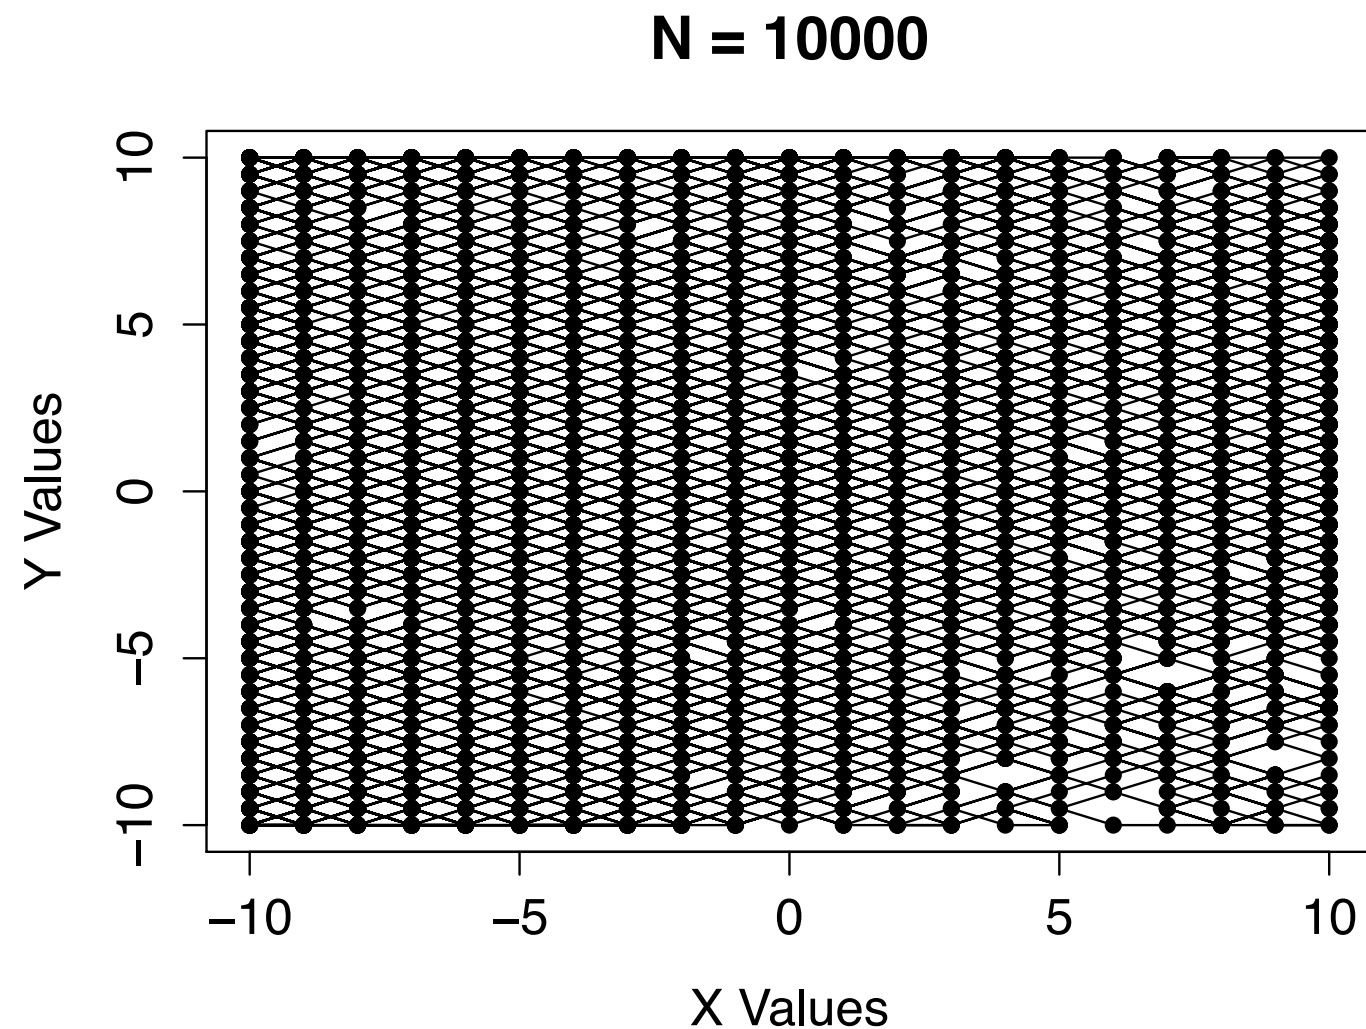
- 1,000 steps (again)



What Is MCMC?

Think of a robot walking through a field

- 10,000 steps



**Your
Turn!**

MCMC Exercise 1

1. Put `basicMCMC.R` file in R's working directory
2. Load it into R

```
source("basicMCMC.R")
```


MCMC Exercise 1

3. Explore how different starting points influence chain. For example, repeat the command below a number of times. (remember the \uparrow functionality). Try other chain lengths too!

```
basicMCMC(nSteps = 100)
```

MCMC Exercise 1

4. Explore how different chain lengths influence chain. Some examples are below, but play around yourself!


```
basicMCMC(nSteps = 100)  
basicMCMC(nSteps = 500)  
basicMCMC(nSteps = 1000)  
basicMCMC(nSteps = 5000)
```

**How Does It Find
Values With Highest
Probabilities?**

How Does It Find High Probabilities?

Clarification

- Instead of a robot walking through a field, is really a **model** moving through **parameter space**



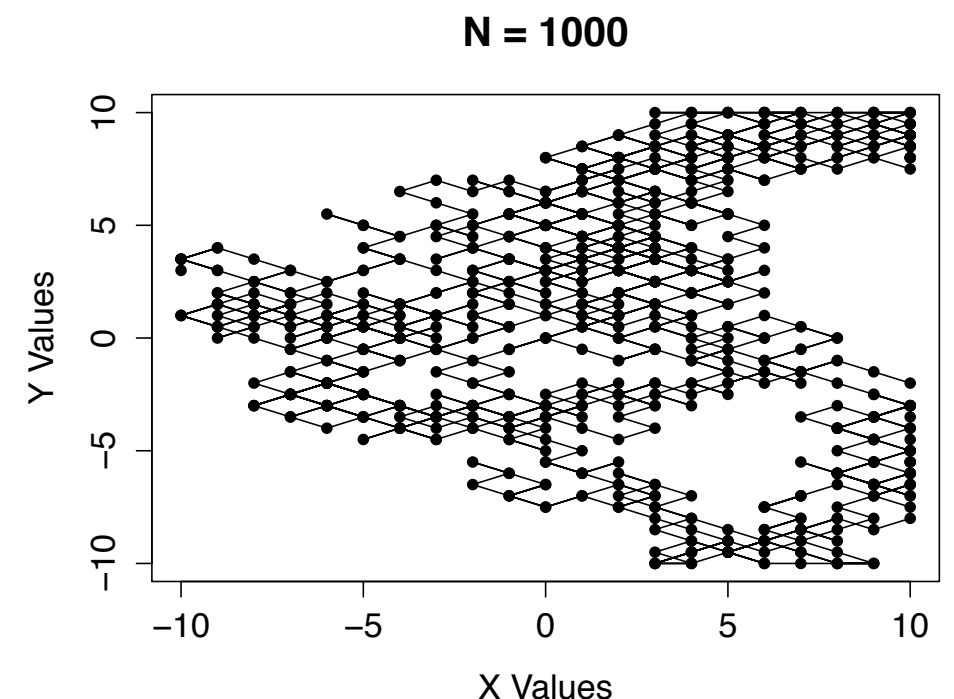
How Does It Find High Probabilities?

Clarification

- Instead of a robot walking through a field, is really a **model** moving through **parameter space**


- Parameter space:
possible x and y values
- Model (simple):
If proposed value is in
possible values (-10 to 10 in this case),
take step

Current parameter space is
flat (no values have higher
likelihood than others)



How Does It Find High Probabilities?

- In real problems, different values will have different likelihoods
- Suppose a simple parameter space with a single peak representing a normal distribution with a mean of 3 and a standard deviation of 2 (in both the x- and y-directions)



How Does It Find High Probabilities?

- The height of the peak for a given value represents the likelihood of that value
- For all probability distributions, this value is calculable (often ugly, but still calculable)
- For the normal distribution, this equation is:

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

x = observed value

μ = mean

σ = s.d.

**Your
Turn!**

Peak-Finding Exercise 1

- Explore how different x values result in different likelihoods


```
dnorm(1, mean = 3, sd = 2)  
dnorm(2, mean = 3, sd = 2)  
etc.
```

Peak-Finding Exercise 1

- Examine the pattern

```
x = seq(from = 0, to = 10, by = 0.5)
y = dnorm(x, mean = 3, sd = 2)
```

```
plot(x, y, type = "b", xlab = "Possible x values", ylab = "Likelihood")
```



How Does It Find High Probabilities?

- Need to set up an **algorithm** that will cause the model to find peaks (if they exist)

Algorithm

A series of instructions for a computer program to follow (like a recipe, but for a computer)


How Does It Find High Probabilities?

- Need to set up an **algorithm** that will cause the model to find peaks (if they exist)
- Example (**algorithm 1**):
 - Step 1: Propose new value (next step)
 - Step 2: Calculate the probability at proposed position
 - Step 3: (a) If proposed probability $>$ probability at current position, take the step
(b) If not, stay in current position

Repeat steps 1-3 many times

How Does It Find High Probabilities?

- **Algorithm 1** will “work”, in that it will climb the first peak it finds and stay there
- However
 - Will get stuck on local peaks (no way to move “downhill”)
 - Won’t explore parameter space very well



How Does It Find High Probabilities?

- Need an algorithm that **preferentially** takes steps to values with higher likelihoods, but has **some potential** to also move to values with lower likelihoods (to walk through valleys)
- **Metropolis-Hastings Algorithm**

How Does It Find High Probabilities?

- **Metropolis-Hastings Algorithm**

- Step 1: Propose new value (next step)
- Step 2: Calculate the probability at proposed position
- Step 3: (a) If proposed probability $>$ probability at current position, take the step
(b) If not,
 - (i) Calculate the ratio of the probabilities at the proposed and current position
 - (ii) Draw a random number between 0 and 1 (inclusive)
 - (iii) If random number is \geq this ratio, take step
 - (iv) If not, don't take step

How Does It Find High Probabilities?

- **Metropolis-Hastings Algorithm**

- Step 1: Propose new value (next
- Step 2: Calculate the probability at proposed position
- Step 3: (a) If proposed probability $>$ probability at current position, take the step
(b) If not,

More readily takes small steps down than big steps, but big steps not impossible

- (i) Calculate the ratio of the probabilities at the proposed and current position
- (ii) Draw a random number between 0 and 1 (inclusive)
- (iii) If random number is \geq this ratio, take step
- (iv) If not, don't take step


How Does It Find High Probabilities?

- Again, the model doesn't have to "know" where peaks are
- If we follow these rules, it should walk around parameter space until it finds a peak, then it should climb the peak and stay there (with some potential to cross valleys)


How Does It Find High Probabilities?

Some examples

N = 100



N = 100



Video

<https://chi-feng.github.io/mcmc-demo/>


How Does It Find High Probabilities?

Heating chains


- Can tweak how stringent conditions are for taking a step to a lower likelihood (downhill)
- Chains that more readily take such steps are called “heated”
 - The more heated a chain is, the more readily it will explore parameter space - good
 - The more heated a chain is, the less time it will stay on top of peaks - bad

Need a balance

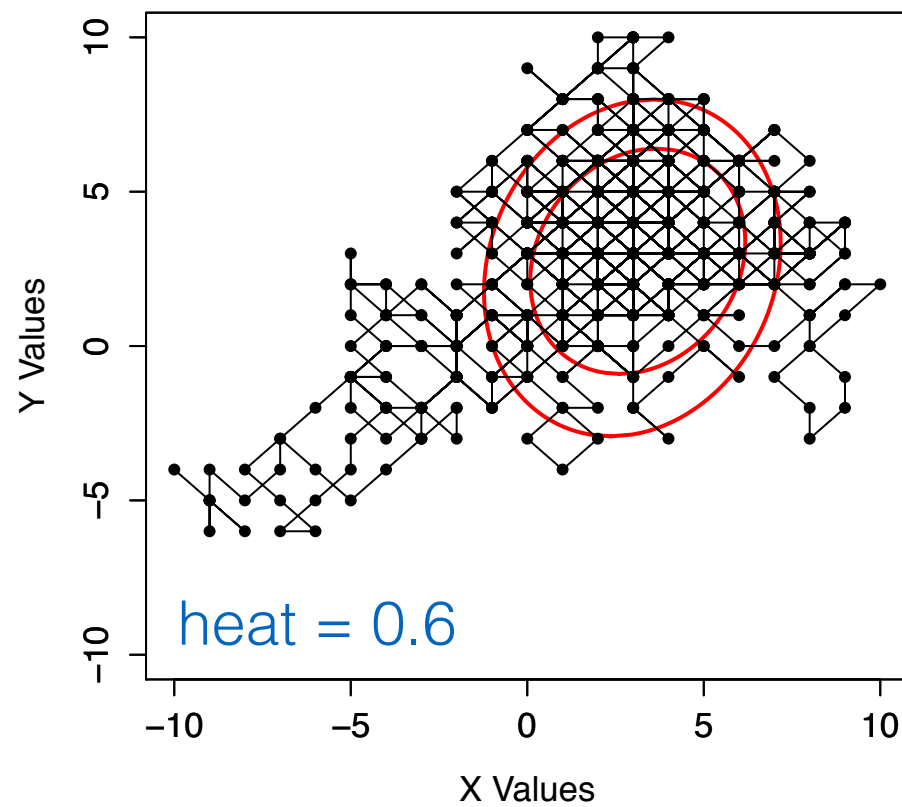
N = 1000




N = 1000



N = 1000



N = 1000



**Your
Turn!**

Peak-Finding Exercise 2

1. Need to install the following libraries
 - a. MASS
 - b. cluster
2. Put `onePeakMCMC.R` file in R's working directory
3. Load it into R

```
source( "onePeakMCMC.R" )
```

Peak-Finding Exercise 2

3. Explore how different starting points influence chain's ability to find peaks. For example, repeat the command below a number of times. (remember the \uparrow functionality). Try other chain lengths too!

```
onePeakMCMC(nSteps = 100, heat = 0, burnin = 0)
```


Peak-Finding Exercise 2

4. Explore how different heating strategies influence how readily the chain steps away from the peak. For example:

```
onePeakMCMC(nSteps = 1000, heat = 0.0, burnin = 0)
onePeakMCMC(nSteps = 1000, heat = 0.2, burnin = 0)
onePeakMCMC(nSteps = 1000, heat = 0.4, burnin = 0)
onePeakMCMC(nSteps = 1000, heat = 0.6, burnin = 0)
onePeakMCMC(nSteps = 1000, heat = 0.8, burnin = 0)
onePeakMCMC(nSteps = 1000, heat = 1.0, burnin = 0)
```

Try some of your own!

**What does this have to
do with Bayesian analysis?**

MCMC & Bayesian Analysis

Like peanut butter & chocolate

- Remember Bayes' Rule

$$P(H|D) = \frac{P(H) \times P(D|H)}{\int P(H) \times P(D|H)}$$

MCMC & Bayesian Analysis

Like peanut butter & chocolate

- Remember Bayes' Rule

$$P(H|D) = \frac{P(H) \times P(D|H)}{\int P(H) \times P(D|H)}$$

Unsolvable for most
problems

MCMC & Bayesian Analysis

Like peanut butter & chocolate

- Remember Bayes' Rule

$$P(H|D) = \frac{P(H) \times P(D|H)}{\int P(H) \times P(D|H)}$$

- If MCMC model is exploring space well:
 - Proportion of time spent at a value represents the posterior probability of that value

MCMC & Bayesian Analysis

Like peanut butter & chocolate

$$P(H|D) = \frac{P(H) \times P(D|H)}{\int P(H) \times P(D|H)}$$

1. Propose new value of (H)

MCMC & Bayesian Analysis

Like peanut butter & chocolate

$$P(H|D) = \frac{P(H) \times P(D|H)}{\int P(H) \times P(D|H)}$$

1. Propose new value of (H)
2. Multiply the prior and likelihood for proposed value

MCMC & Bayesian Analysis

Like peanut butter & chocolate

$$P(H|D) = \frac{P(H) \times P(D|H)}{\int P(H) \times P(D|H)}$$

1. Propose new value of (H)
2. Multiply the prior and likelihood for proposed value
 - a. If this product is \geq current value, take step
 - b. If this product is \leq current value, follow some rules to determine whether or not to take step

MCMC & Bayesian Analysis

Like peanut butter & chocolate

$$P(H|D) = \frac{P(H) \times P(D|H)}{\int P(H) \times P(D|H)}$$

1. Propose new value of (H)
2. Multiply the prior and likelihood for proposed value
 - a. If this product is \geq current value, take step
 - b. If this product is \leq current value, follow some rules to determine whether or not to take step
3. If MCMC model is exploring space well, **proportion** of time spent at a specific value of H will represent its posterior probability

MCMC & Bayesian Analysis

Like peanut butter & chocolate

- The benefits of this cannot be overstated
- Allows easy computation of very difficult problems
 - Problems not solvable (or even fathomable) before MCMC & good computing power
- Only downside is that MCMC is computationally “expensive”

MCMC & Bayesian Analysis

Like peanut butter & chocolate


- Can easily accommodate problems dealing with 10s or even 100s of parameters
 - As long as likelihood for each can be calculated...
 - ... and model is set up properly

Burn-in


Burn-in

- To be useful in Bayesian analyses, proportion of time spent at a value should be reflective of that value's probability
- But, remember that the model wanders aimlessly, and is heavily biased by where it starts, until it actually finds a peak
- Need to ignore these steps when making estimation
 - Number of steps ignored is called the "burn-in"
 - No way to know ahead of time, but better to ignore too many than too few


N = 1000




N = 1000



N = 1000



N = 1000



**Your
Turn!**

Burn-in Exercise

1. Explore how different burn-in lengths fit with different chain lengths and heating values.

```
onePeakMCMC(nSteps = 1000, heat = 0.5, burnin = 10)
onePeakMCMC(nSteps = 1000, heat = 0.5, burnin = 50)
onePeakMCMC(nSteps = 1000, heat = 0.5, burnin = 100)
etc.
```

Try some of your own!

Burn-in


- Not really possible to have burn-in “too long” (but you need enough steps of data collection to obtain good estimates)
- A burn-in that is too short can cause real problems

We'll go over some ways to examine this later

Running Multiple Chains

Running Multiple Chains

- We've been talking about very simple examples
- Parameter space for "real" problems are likely quite complex (and unknowingly so)



Running Multiple Chains


- Often more efficient to run multiple chains
 - Each starting in a different place
 - Each with a different heating value

Will explore parameter space more efficiently than
a single chain

Running Multiple Chains

- Have them swap periodically
 - Have one main chain, and the rest “workers”
 - Every so often, compare the probability of positions of each chain
 - If a worker chain has a higher probability than the main chain, switch parameter values between chains

Allows main chain to “jump”
to peaks that other chains have
found



Hamiltonian Monte Carlo




Hamiltonian Monte Carlo

- A physics simulation
- Represent parameter state as a particle in n -dimensional space
- Flick particle around frictionless $-\log$ -posterior
- Record positions


Hamiltonian Monte Carlo

```
x = seq(from = 0, to = 1, by = 0.01)
y = dnorm(x, mean = 0.5, sd = 0.2)
plot(y ~ x, type = "l")
```




Hamiltonian Monte Carlo

```
y2 = -log(y)  
plot(y2 ~ x, type = "l")
```



Hamiltonian Monte Carlo



- Position
 - Present parameter values
- Speed (momentum of parameter values)
 - Fast when parameter values have low probability (high $-\log\text{-prob}$)
 - Slow when parameter values have high probability (low $-\log\text{-prob}$)
- Samples of position are samples from posterior

Video

(Naïve HMC)

<https://chi-feng.github.io/mcmc-demo/>

Hamiltonian Monte Carlo

- When do you take estimate (stop the particle)?
- NUTS: No U-turn sampler
 - Start trajectory in both directions
 - Run until it starts to turn around



Video

<https://chi-feng.github.io/mcmc-demo/>

Hamiltonian Monte Carlo

- Does not have *burn-in*
- Has *warmup* instead
 - Stan is optimizing model behaviour for your model (step size, mass of particle, etc.)

Questions?