

Hierarchical Modeling

Tim Frasier

General Idea


- One benefit of a Bayesian approach is that you have full control of your models and what assumptions they make about the data
- Hierarchical modelling provides a more efficient use of the data, where you can explicitly model how different estimates are related to one another
- Will start simple now, but get more complicated later

Rationale

- Current models have amnesia:
 - Each category has completely new prior distribution
 - As if model has never seen data like it before
- Is this best use of data?
 - Surely different categories of the same variable provide *some* information about the other categories
 - Height coefficient of one sex must provide *some* information about the coefficient for the other sex (can't differ by 0, or by 1,000,000)




Example

- Programming a robot to estimate waiting times at coffee shops
 - Tim Hortons
 - Starbucks
 - Second Cup
- When switching from Tim Hortons to Starbucks, like its never been to a coffee shop before






```
// Likelihood
for (i in 1:N) {
  mu[i] = b1[x[i]];
  y[i] ~ normal(mu[i], sigma[x[i]]);
}

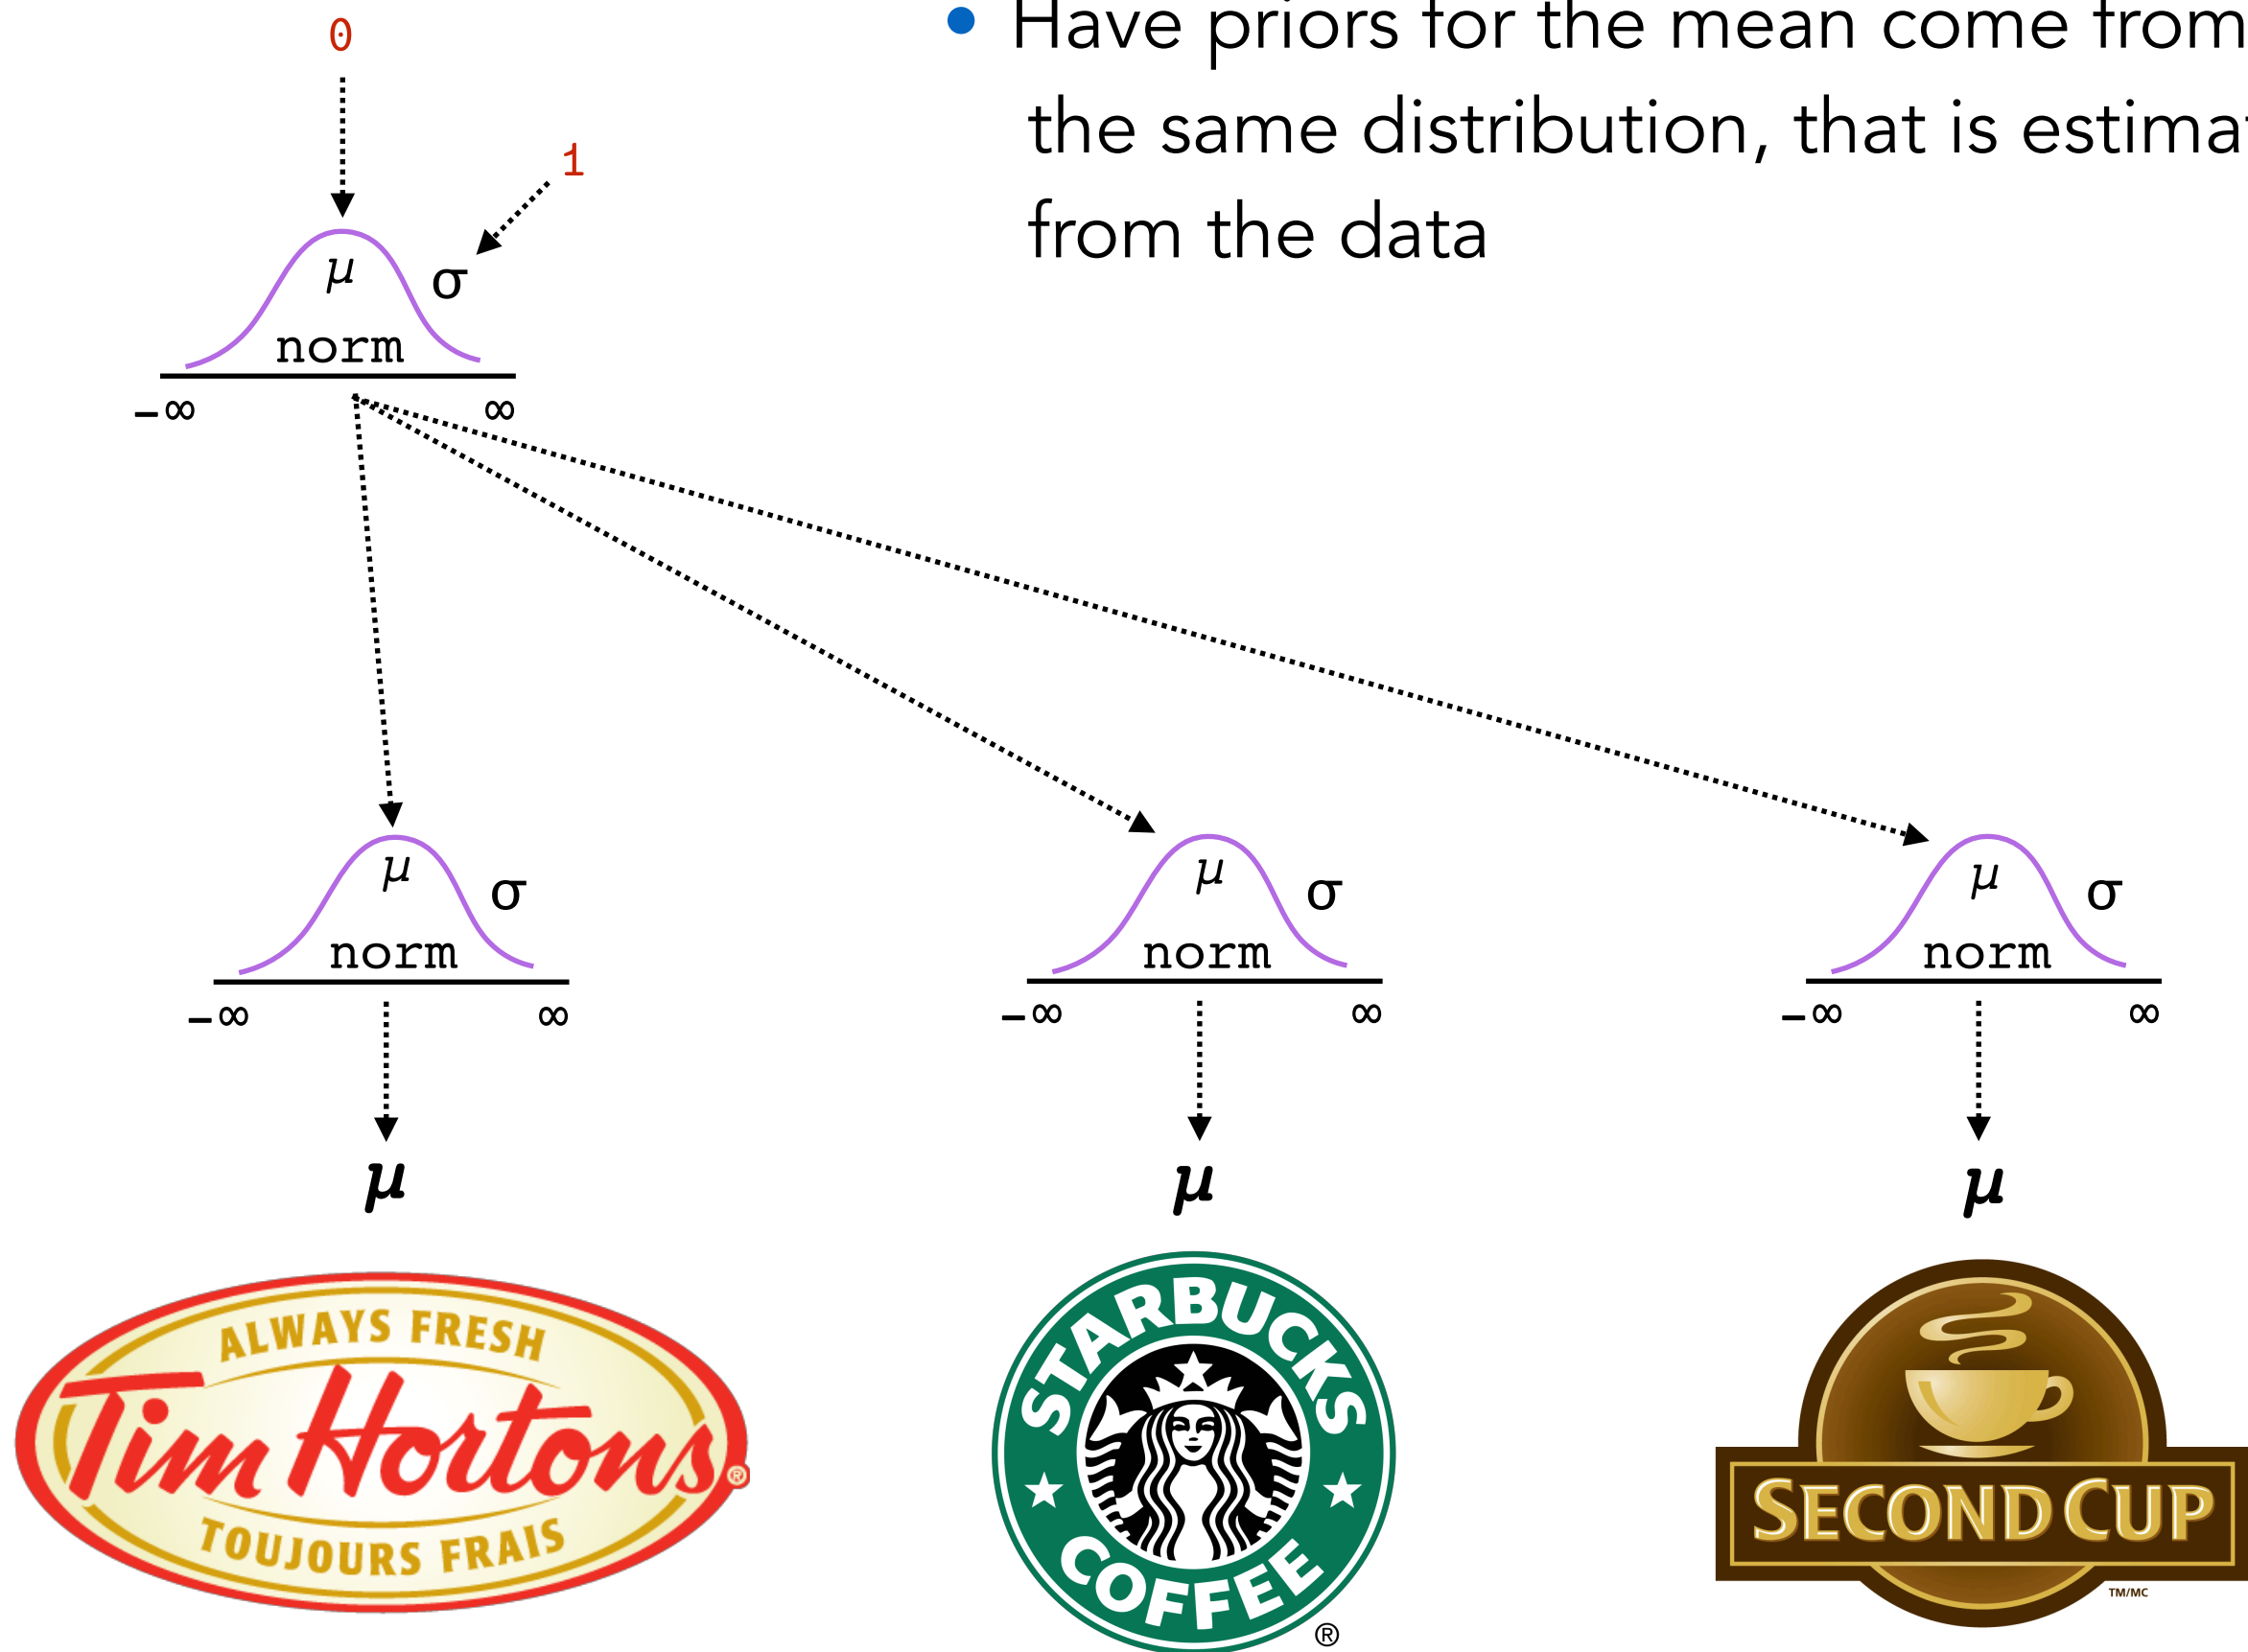
// Priors
for (j in 1:nxLevels) {
  b1[j] ~ normal(0, 1);
  sigma[j] ~ cauchy(1, 1);
}
```



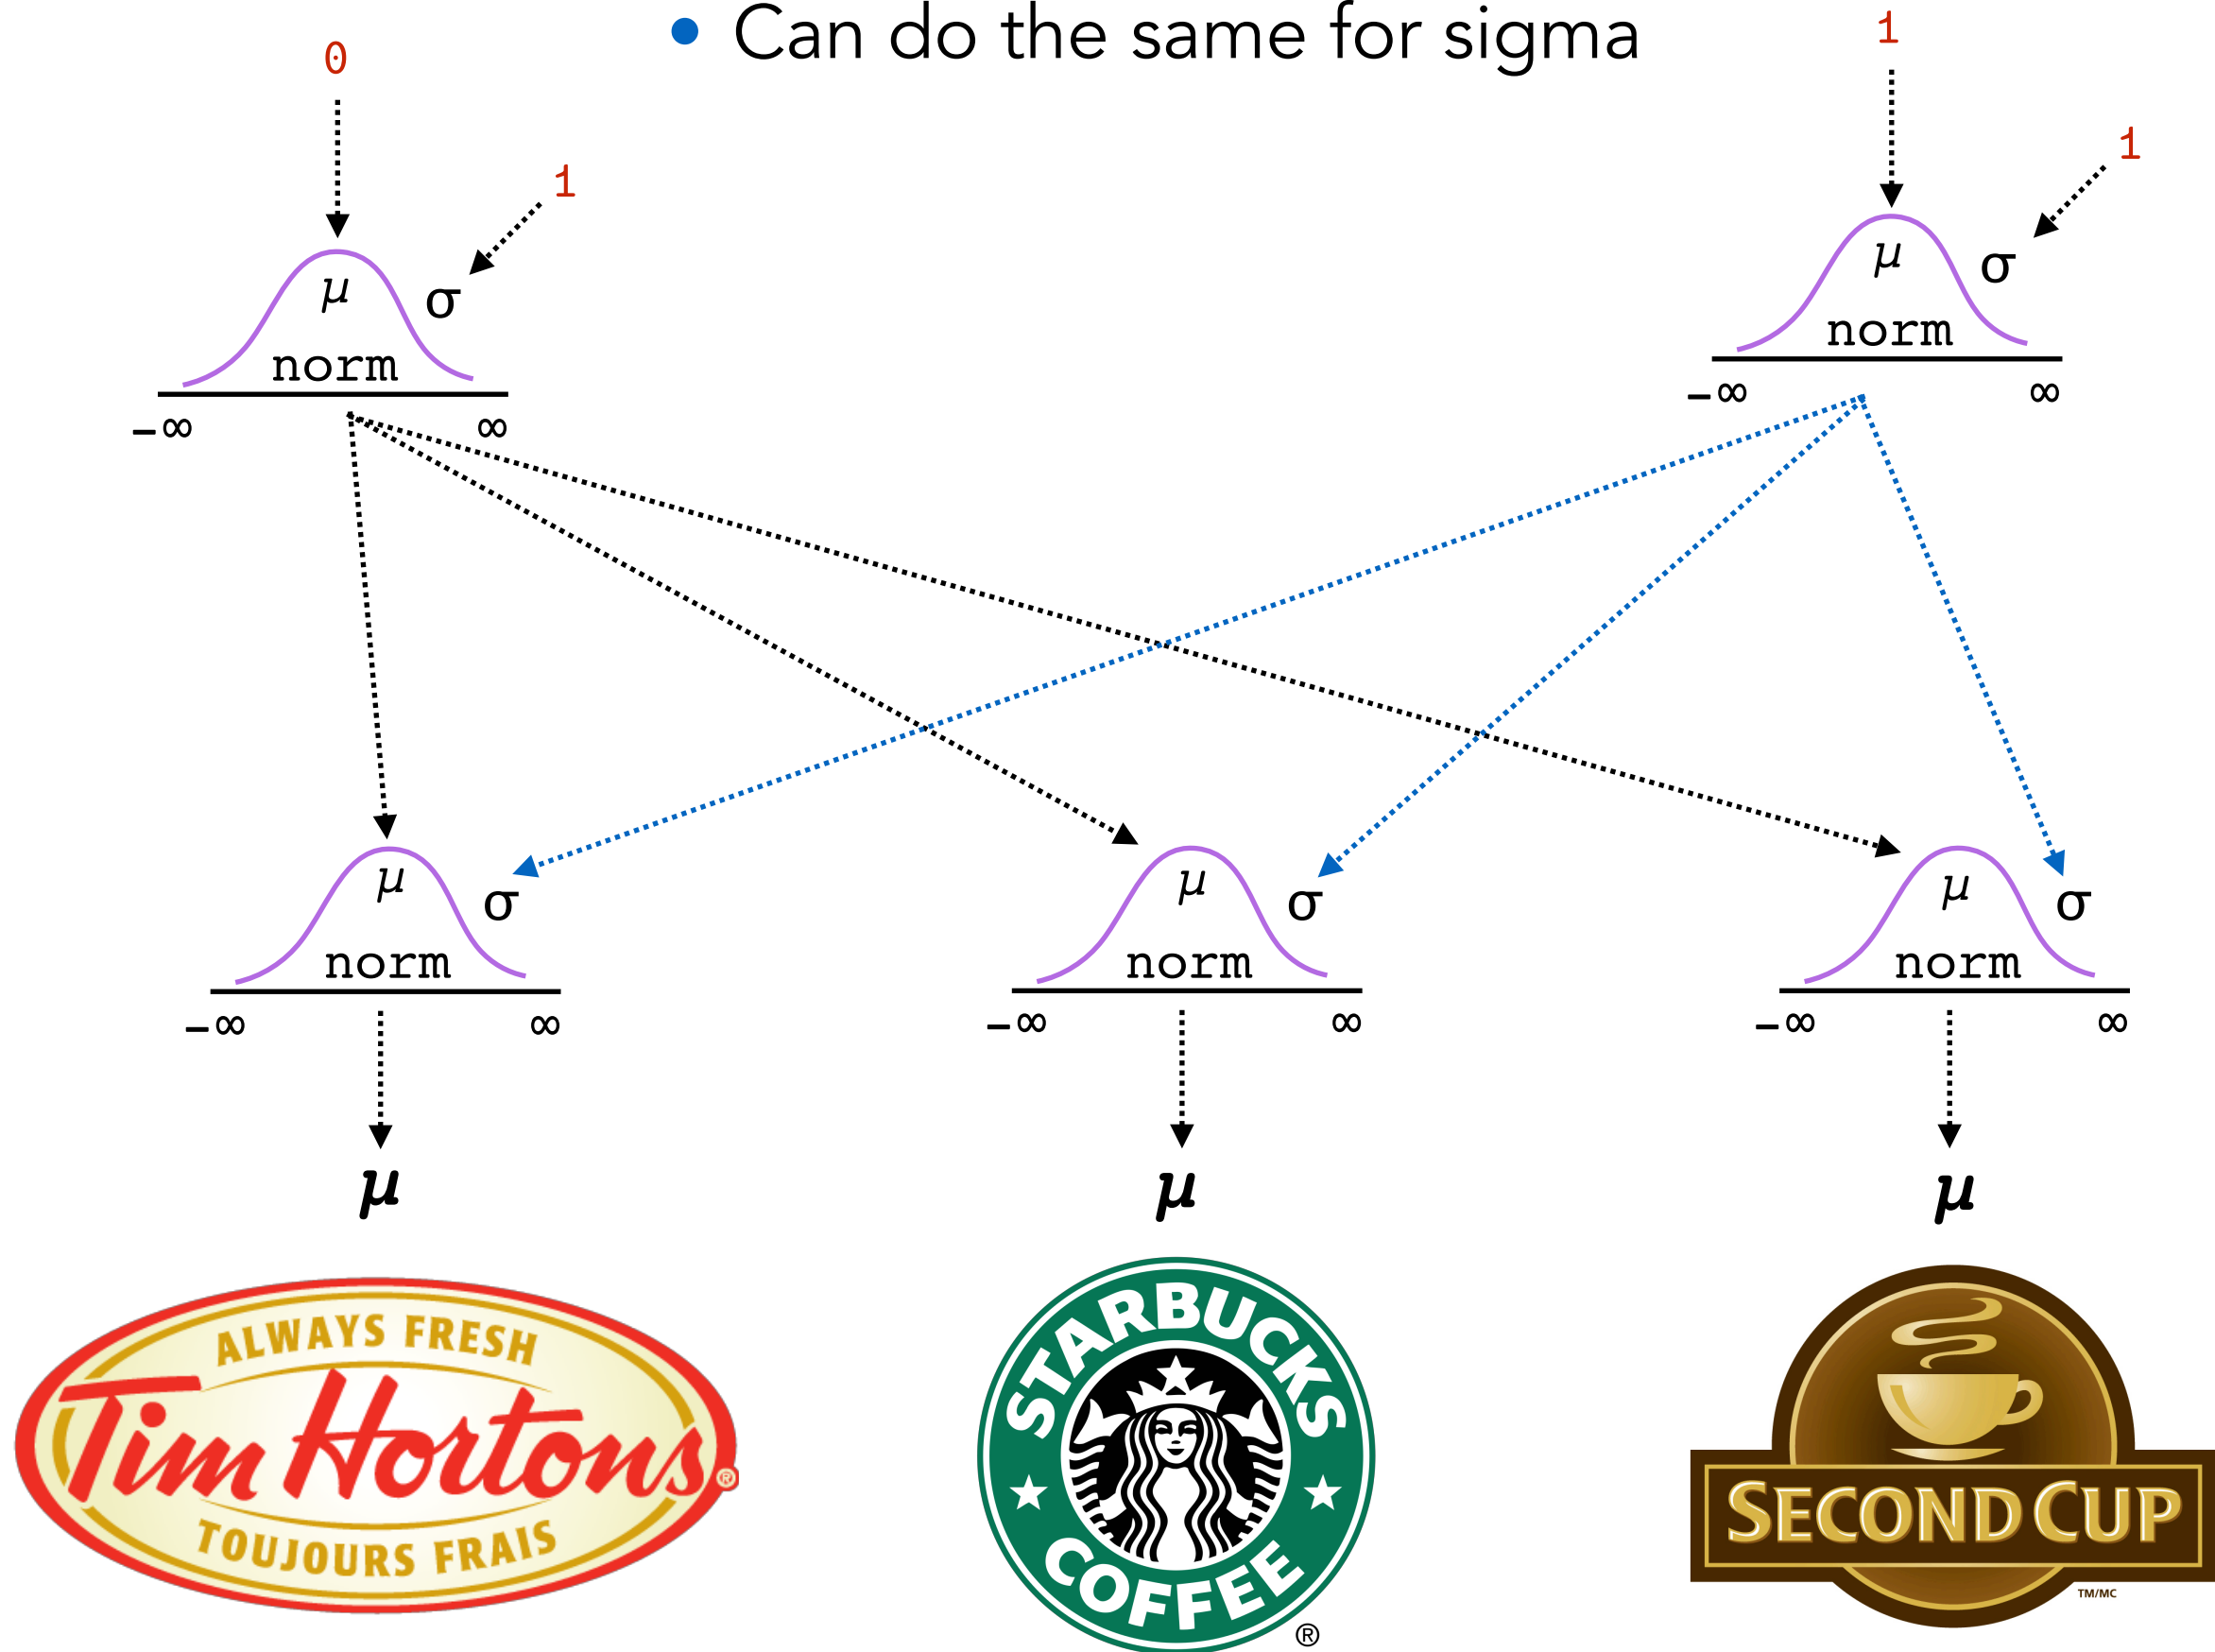
- Can have each data point be part of the estimate for that category, while also informing the overall pattern
- Sort of like having the posterior from one category be the prior for the next



- Have priors for the mean come from the same distribution, that is estimated from the data

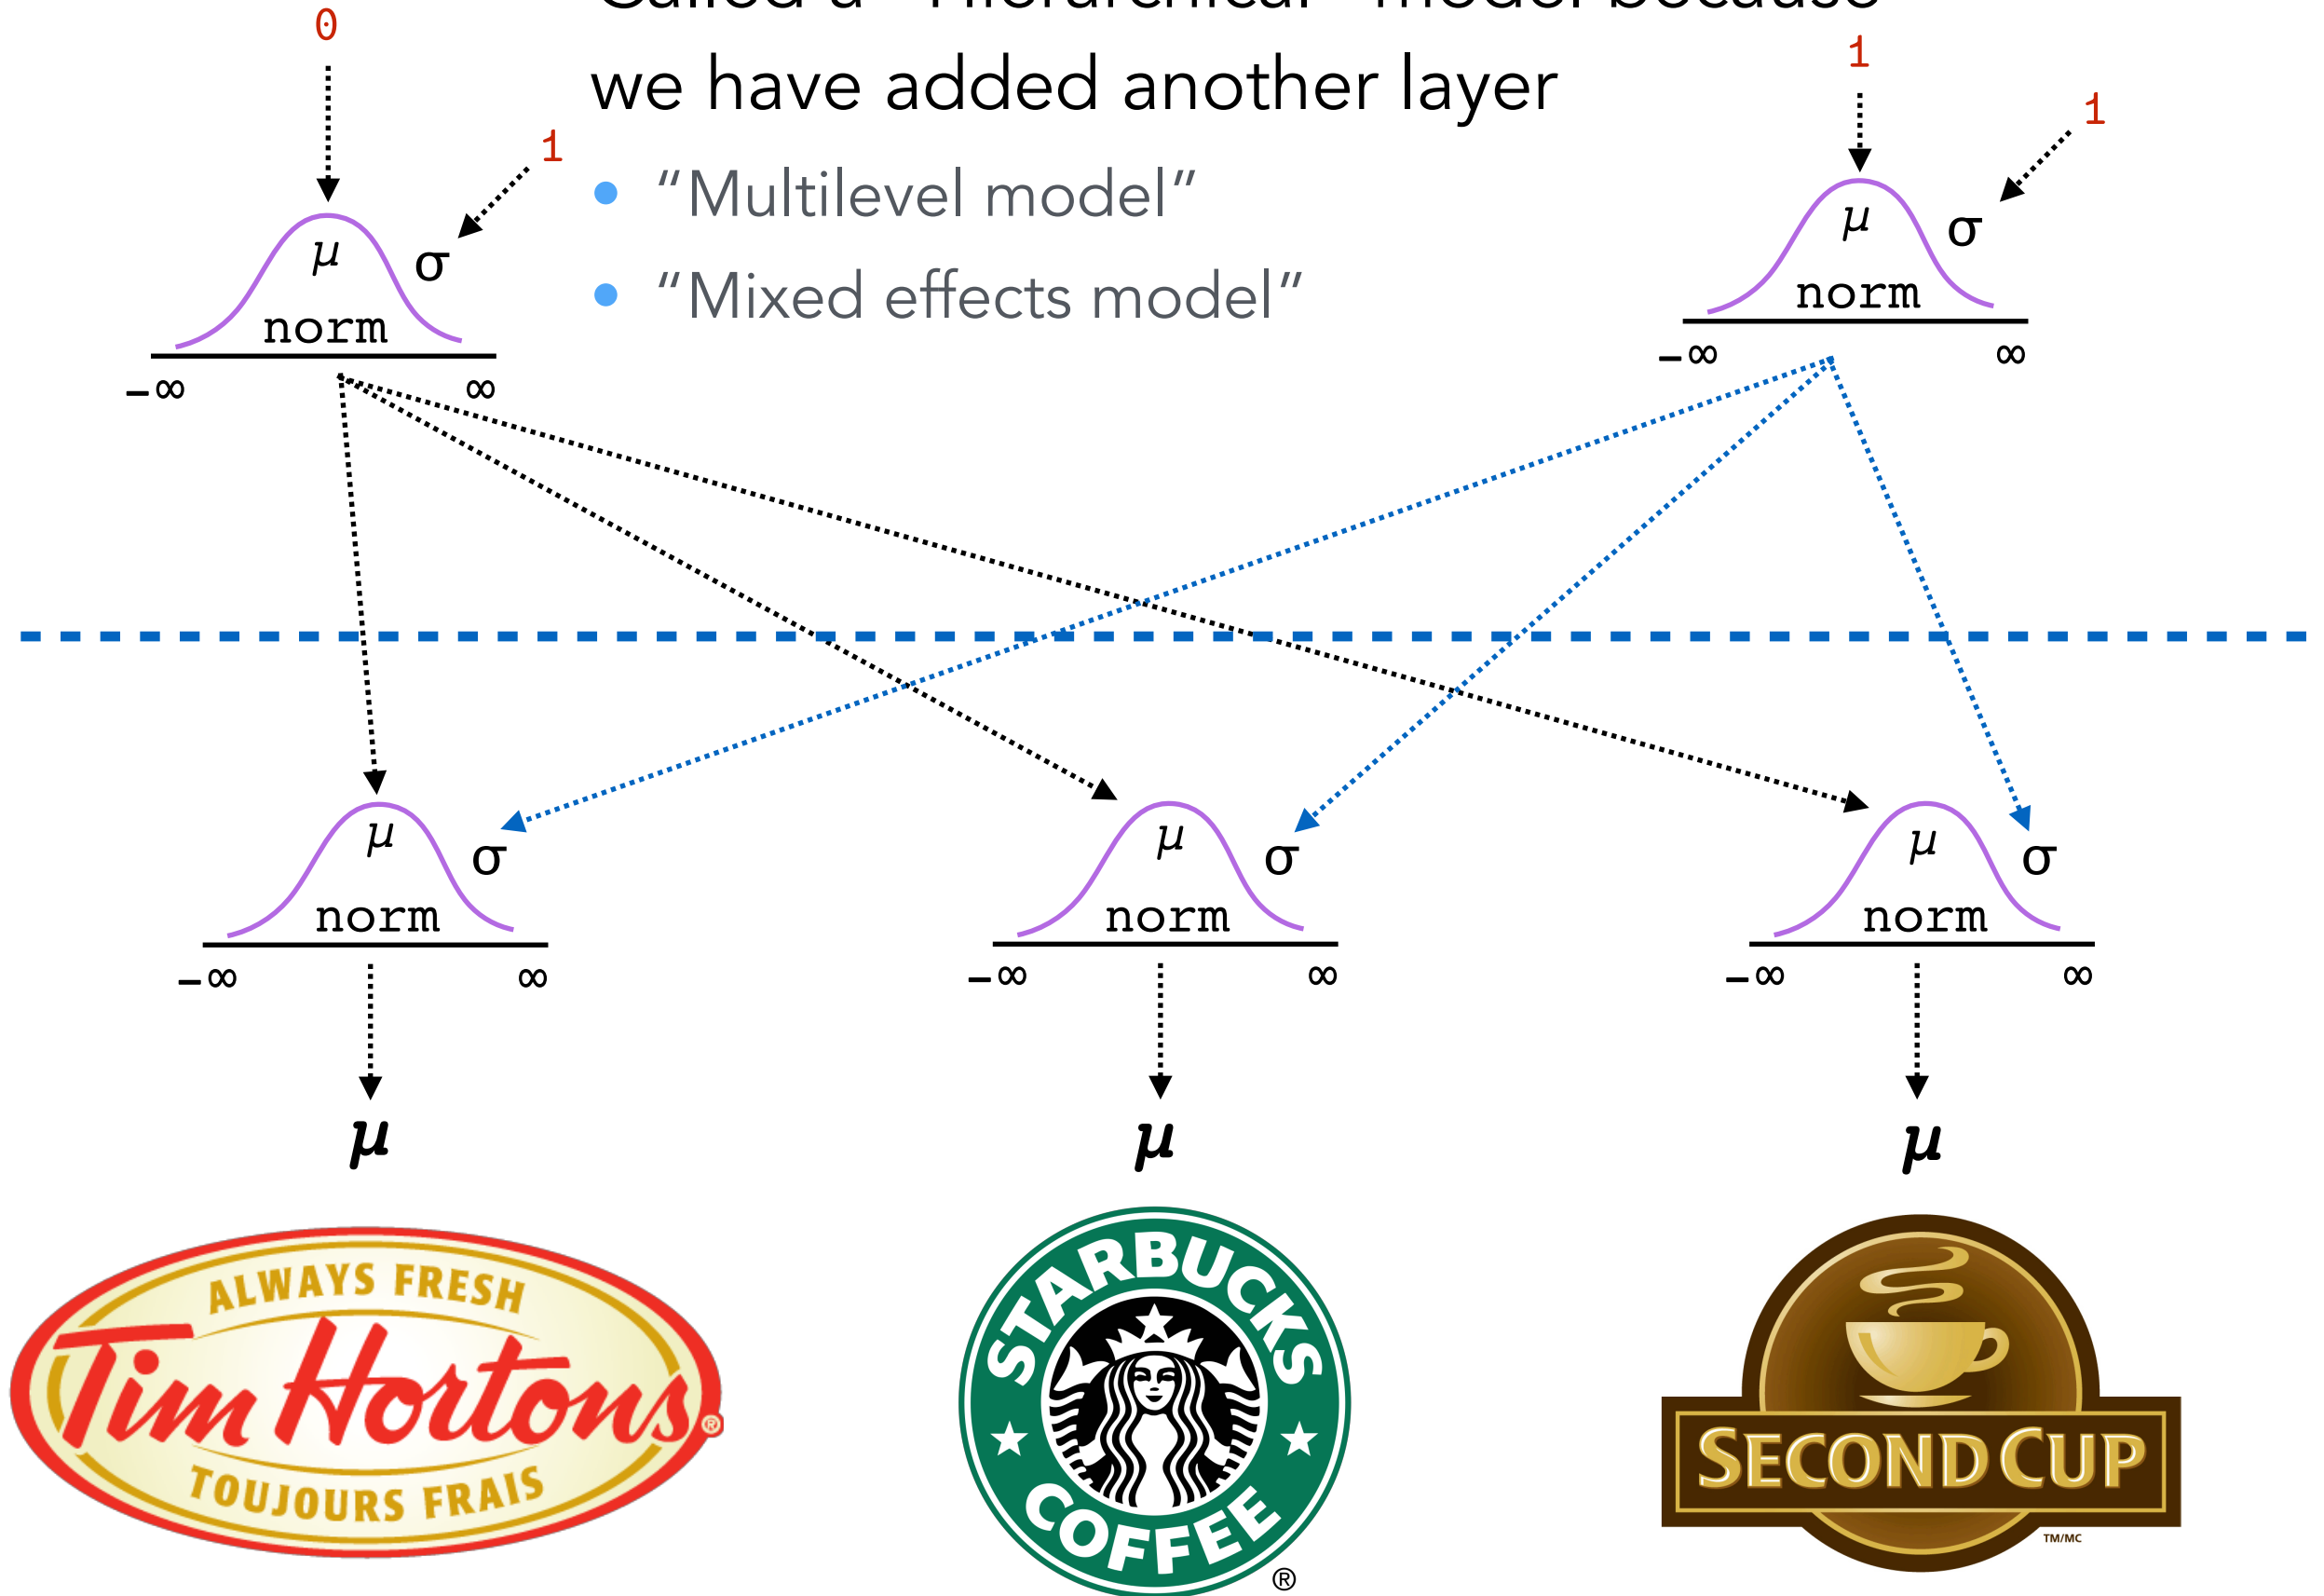


- Can do the same for sigma

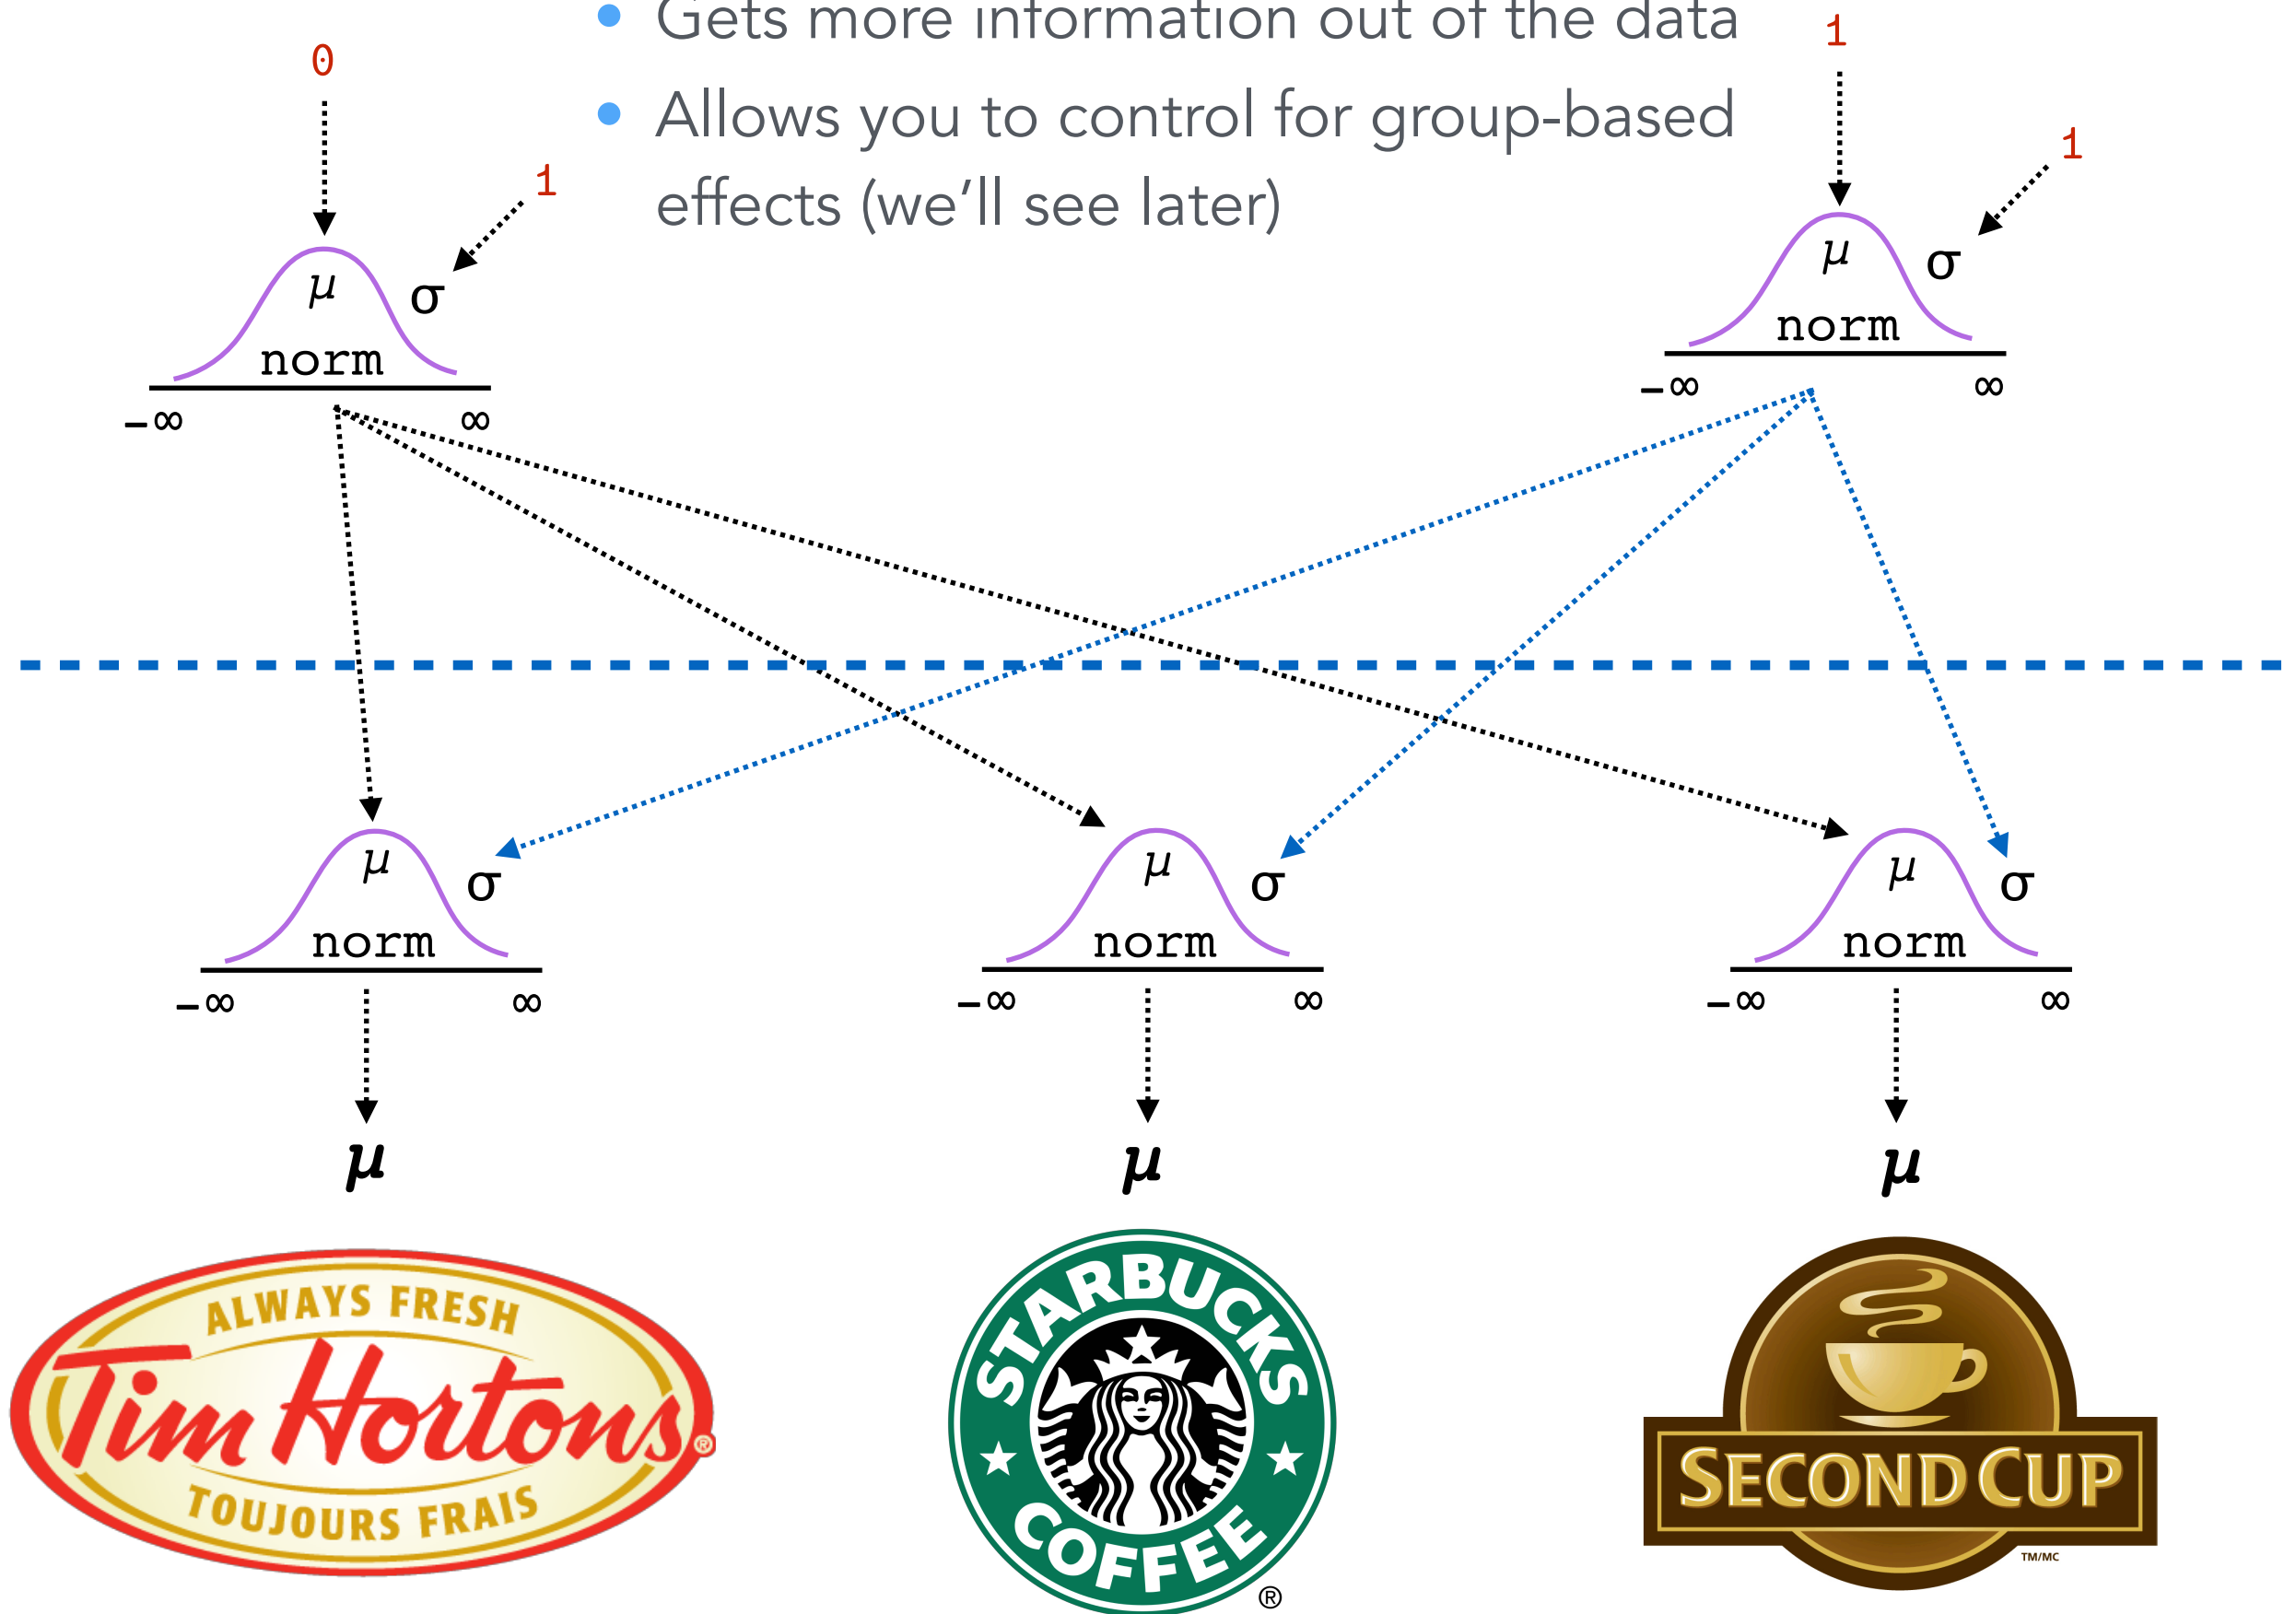


- Called a “Hierarchical” model because we have added another layer

- “Multilevel model”
- “Mixed effects model”



- Gets more information out of the data
- Allows you to control for group-based effects (we'll see later)



The Data

Data

- `coffee.csv`



30 trips, true average = 15 minutes



5 trips, true average = 5 minutes




15 trips, true average = 10 minutes

coffee			
	trip	shop	times
1	1	Tim Hortons	10.0991578
2	2	Tim Hortons	4.1515614
3	3	Tim Hortons	12.4475285
4	4	Tim Hortons	9.6105939
5	5	Tim Hortons	13.5664162
6	6	Tim Hortons	7.0864281
7	7	Tim Hortons	16.6772393
8	8	Tim Hortons	7.8437422
9	9	Tim Hortons	11.8697399
10	10	Tim Hortons	7.5585060
11	11	Tim Hortons	19.1699169

Data


```
coffee = read.table("coffee.csv", header = TRUE, sep = ",")

ggplot(coffee) +
  theme_bw() +
  geom_boxplot(aes(x = shop, y = time), fill = "dodgerblue4", alpha = 0.6)
```




Data

```
ggplot(coffee) +  
  theme_bw() +  
  geom_jitter(aes(x = shop, y = time), height = 0, width = 0.3, alpha = 0.6)
```



Data



Load Libraries & Functions

```
library(rstan)  
library(ggplot2)  
source("plotPost.R")
```


Prepare the Data

Standardize metric (y) data

```
y = coffee$time  
yMean = mean(y)  
ySD = sd(y)  
zy = (y - yMean) / ySD  
  
N = length(zy)
```

Prepare the Data

Organize the nominal (x) data

```
x = as.numeric(coffee$shop)

xNames = levels(as.factor(coffee$shop))

nxLevels = length(unique(coffee$shop))
```

Create a data list for Stan

```
dataList = list(  
  y = zy,  
  x = x,  
  N = N,  
  nxLevels = nxLevels  
)
```

Build the Model

- The **data block** stays the same

```
data {  
  int N;           // Sample size  
  int nxLevels;    // Number of shop types (levels in our nominal variable)  
  vector[N] y;     // Vector of wait times  
  int x[N];        // Vector of indicators of which level each sample is in  
}
```

Build the Model

- Let's look at the **model block** next
 - Definitions and likelihood remain the same

```
model {  
  // Definitions  
  vector[N] mu;  
  
  // Likelihood  
  for (i in 1:N) {  
    mu[i] = b1[x[i]];  
    y[i] ~ normal(mu[i], sigma[x[i]]);  
  
    ...  
  }  
}
```

Build the Model

- Let's look at the **model block** next
 - Priors are where things get interesting

```
...  
// Priors  
  for (j in 1:nxLevels) {  
    b1[j]      ~ normal(shopMean, shopMeanSD);  
    sigma[j] ~ cauchy(1, 1);  
  }  
  
  // Hyperpriors  
  shopMean      ~ normal(0, 1);  
  shopMeanSD ~ normal(1, 1);  
}
```

- Old model

```
...  
// Priors  
  for (j in 1:nxLevels) {  
    b1[j] ~ normal(0, 1);  
    sigma[j] ~ cauchy(1, 1);  
  }
```

- New model

```
...  
// Priors  
  for (j in 1:nxLevels) {  
    b1[j] ~ normal(shopMean, shopMeanSD);  
    sigma[j] ~ cauchy(1, 1);  
  }  
  
  // Hyperpriors  
  shopMean ~ normal(0, 1);  
  shopMeanSD ~ normal(1, 1);  
}
```

- Old model

```
...  
// Priors  
for (j in 1:nxLevels) {  
  b1[j] ~ normal(0, 1);  
  sigma[j] ~ cauchy(1, 1);  
}
```

Going to estimate a normal distribution from which **all** shop effects arise (will be informed by the distribution of each). Thus we need to estimate the mean and s.d. of this overall distribution.

- New model

```
...  
// Priors  
for (j in 1:nxLevels) {  
  b1[j] ~ normal(shopMean, shopMeanSD);  
  sigma[j] ~ cauchy(1, 1);  
}  
  
// Hyperpriors  
shopMean ~ normal(0, 1);  
shopMeanSD ~ normal(1, 1);  
}
```


- Old model

```
...  
// Priors  
for (j in 1:nxLevels) {  
  b1[j] ~ normal(0, 1);  
  sigma[j] ~ cauchy(1, 1);  
}
```

Going to estimate a normal distribution from which **all** shop effects arise (will be informed by the distribution of each). Thus we need to estimate the mean and s.d. of this overall distribution.

- New model

```
...  
// Priors  
for (j in 1:nxLevels) {  
  b1[j] ~ normal(shopMean, shopMeanSD);  
  sigma[j] ~ cauchy(1, 1);  
}
```

```
// Hyperpriors  
shopMean ~ normal(0, 1);  
shopMeanSD ~ normal(1, 1);  
}
```

Because we are estimating the mean and s.d. of this distribution, they need priors (yes...the priors need priors)

Build the Model

- Go back to the **parameters block**

```
parameters {  
  real b1[nxLevels];           // Coefficients for effects of being in each level  
  real<lower=0> sigma[nxLevels]; // Coefficients for sd of being in each level  
  real shopMean;               // Mean wait time across all shops  
  real shopMeanSD;             // sd for mean wait time across all shops  
}
```

Build the Model

- Lastly, the **generated quantities block** (stays the same)

```
generated quantities {  
  // Definitions  
  vector[N] y_pred;  
  
  for (i in 1:N) {  
    y_pred[i] = normal_rng(b1[x[i]], sigma[x[i]]);  
  }  
}
```

Run the Model

- Remember to include all parameters you are interested in!!!!

```
stanFit <- stan(file = "model.stan",  
               data = dataList,  
               pars = c("b1", "sigma", "shopMean", "shopMeanSD", "y_pred"),  
               warmup = 2000,  
               iter = 7000,  
               chains = 3)
```

Check MCMC Performance

Rhat


```
print(stanFit)
```

[illegible]

Check MCMC Performance

traceplots


```
stan_trace(stanFit, pars = "b1", inc_warmup = TRUE)
```



Check MCMC Performance

traceplots

```
stan_trace(stanFit, pars = "sigma", inc_warmup = TRUE)
```



Evaluate Results

Extract the Data

```
mcmcChains = as.data.frame(stanFit)

zshopMean = mcmcChains$shopMean
zshopMeanSD = mcmcChains$shopMeanSD

chainLength = length(mcmcChains[, 1])

zsigma = matrix(0, ncol = nxLevels, nrow = chainLength)
for (i in 1:nxLevels) {
  zsigma[, i] = mcmcChains[, paste("sigma[", i, "]", sep = "")]
}
```


Evaluate Results

Extract the Data

```
zb1 = matrix(0, ncol = nxLevels, nrow = chainLength)
for (i in 1:nxLevels) {
  zb1[, i] = mcmcChains[, paste("b1[, i, ]", sep = "")]
}

zypred = matrix(0, ncol = N, nrow = chainLength)
for (i in 1:N) {
  zypred[, i] = mcmcChains[, paste("y_pred[, i, ]", sep = "")]
}
```

Evaluate Results

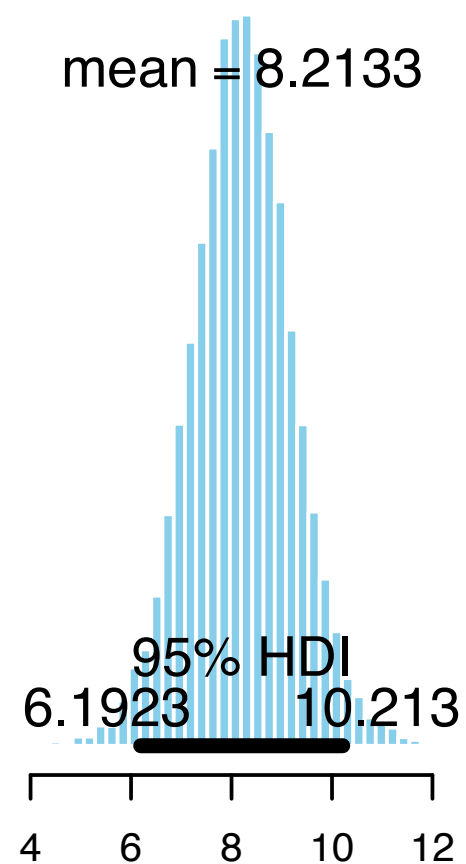
Convert back to original scale

```
sigma = zsigma * ySD  
b1 = (zb1 * ySD) + yMean  
ypred = (zypred * ySD) + yMean  
shopMean = (zshopMean * ySD) + yMean  
shopMeanSD = zshopMeanSD * ySD
```

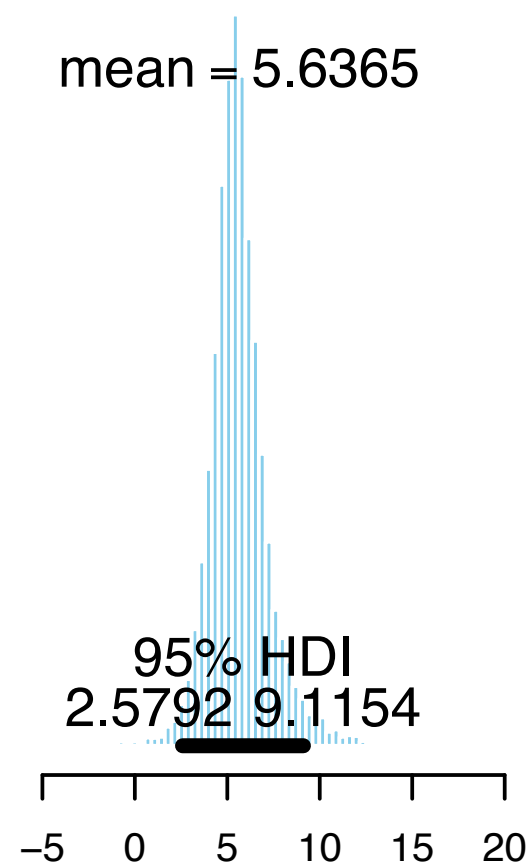
Evaluate Results

Plot estimates

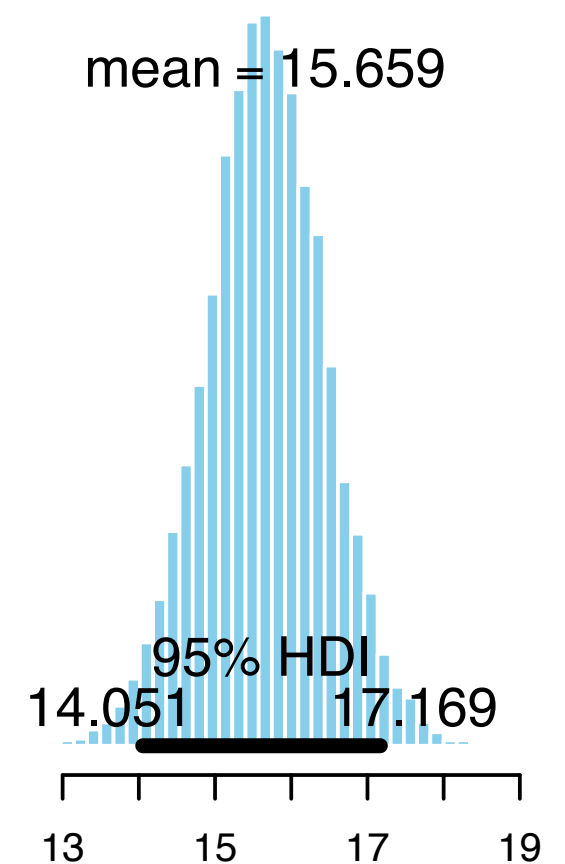
```
par(mfrow = c(1, 3))  
histInfo = plotPost(b1[, 1], xlab = "Wait time: Second Cup")  
histInfo = plotPost(b1[, 2], xlab = "Wait time: Starbucks")  
histInfo = plotPost(b1[, 3], xlab = "Wait time: Tim Hortons")
```



Wait time: Second Cup



Wait time: Starbucks






Wait time: Tim Hortons

Evaluate Results

Plot estimates



```
histInfo = plotPost(b1[, 1], xlab = "Wait time: Second Cup")  
histInfo = plotPost(b1[, 2], xlab = "Wait time: Starbucks")  
histInfo = plotPost(b1[, 3], xlab = "Wait time: Tim Hortons")
```



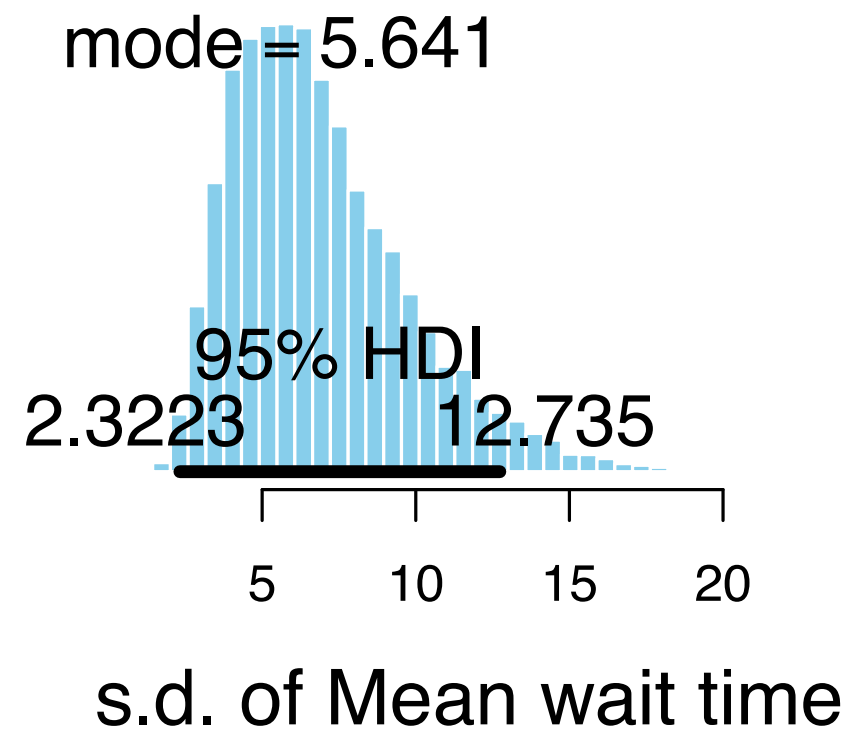

Evaluate Results

Plot estimates

```
par(mfrow = c(1, 2))  
histInfo = plotPost(shopMean, xlab = "Mean wait time")  
histInfo = plotPost(shopMeanSD, xlab = "s.d. of Mean wait time", showMode = TRUE)
```




- What does a normal distribution with a mean of 10.671, and a s.d. of 5.641 look like?



- What does a normal distribution with a mean of 10.671, and a s.d. of 5.641 look like?


```
x = seq(from = 0, to = 40, length = 200)
y = dnorm(x, mean = 10.671, sd = 5.641)

plot(y ~ x, type = "l", ylim = c(0, 0.15), lwd = 3, col = rgb(0, 0, 1, 0.6))
```



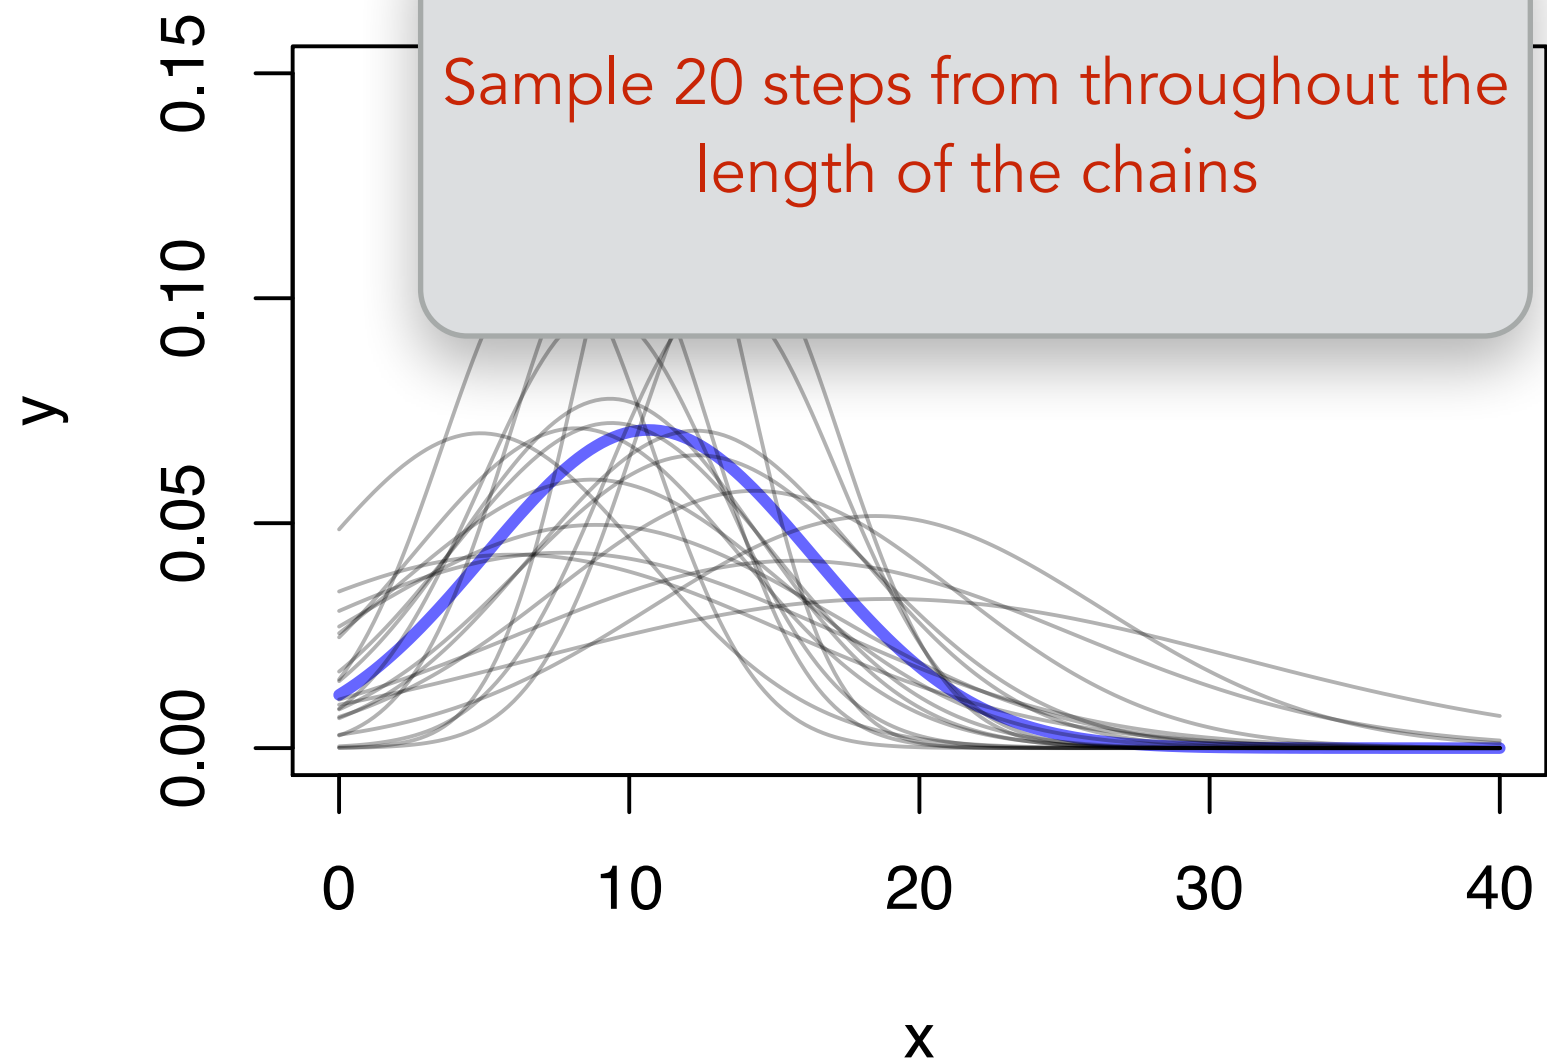
- What about the variation!?

```
samps1 = seq(from = 1, to = length(shopMean), length = 20)
samps2 = round(samps1)
for (i in samps2) {
  y = dnorm(x, mean = shopMean[i], sd = shopMeanSD[i])
  lines(x, y, col = rgb(0, 0, 0, 0.3))
}
```



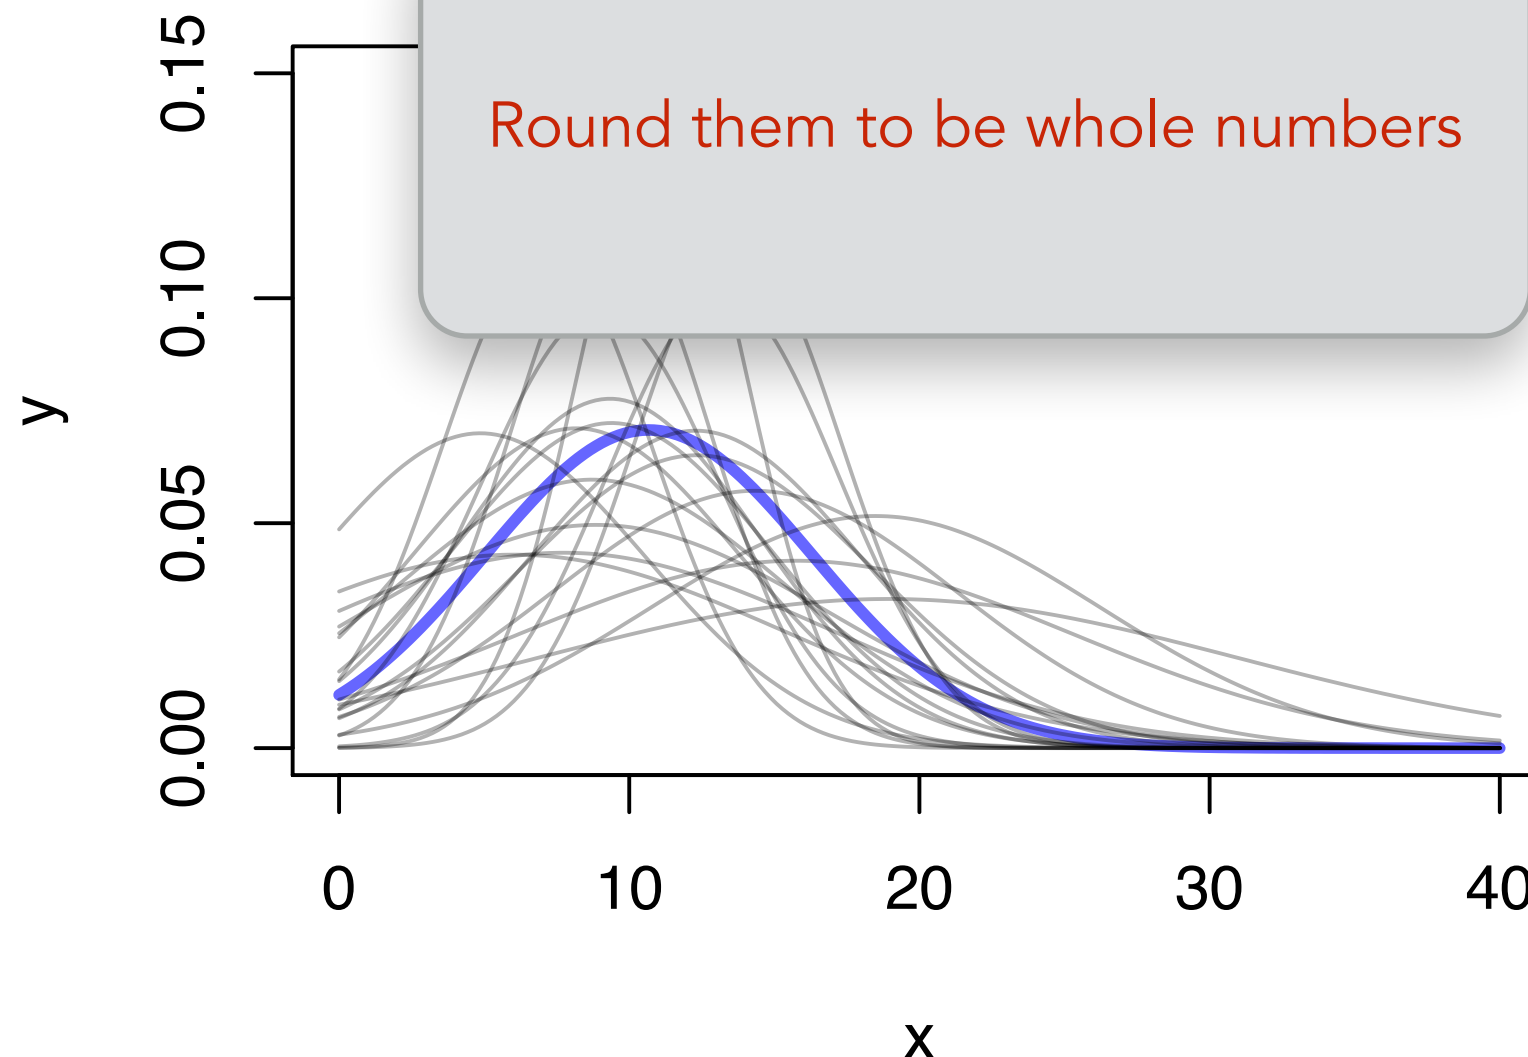
- What about the variation!?

```
samps1 = seq(from = 1, to = length(shopMean), length = 20)
samps2 = round(samps1)
for (i in samps2) {
  y = dnorm(x, mean = shopMean[i], sd = shopMeanSD[i])
  lines(x, y, col = rgb(0, 0, 0, 0.3))
}
```



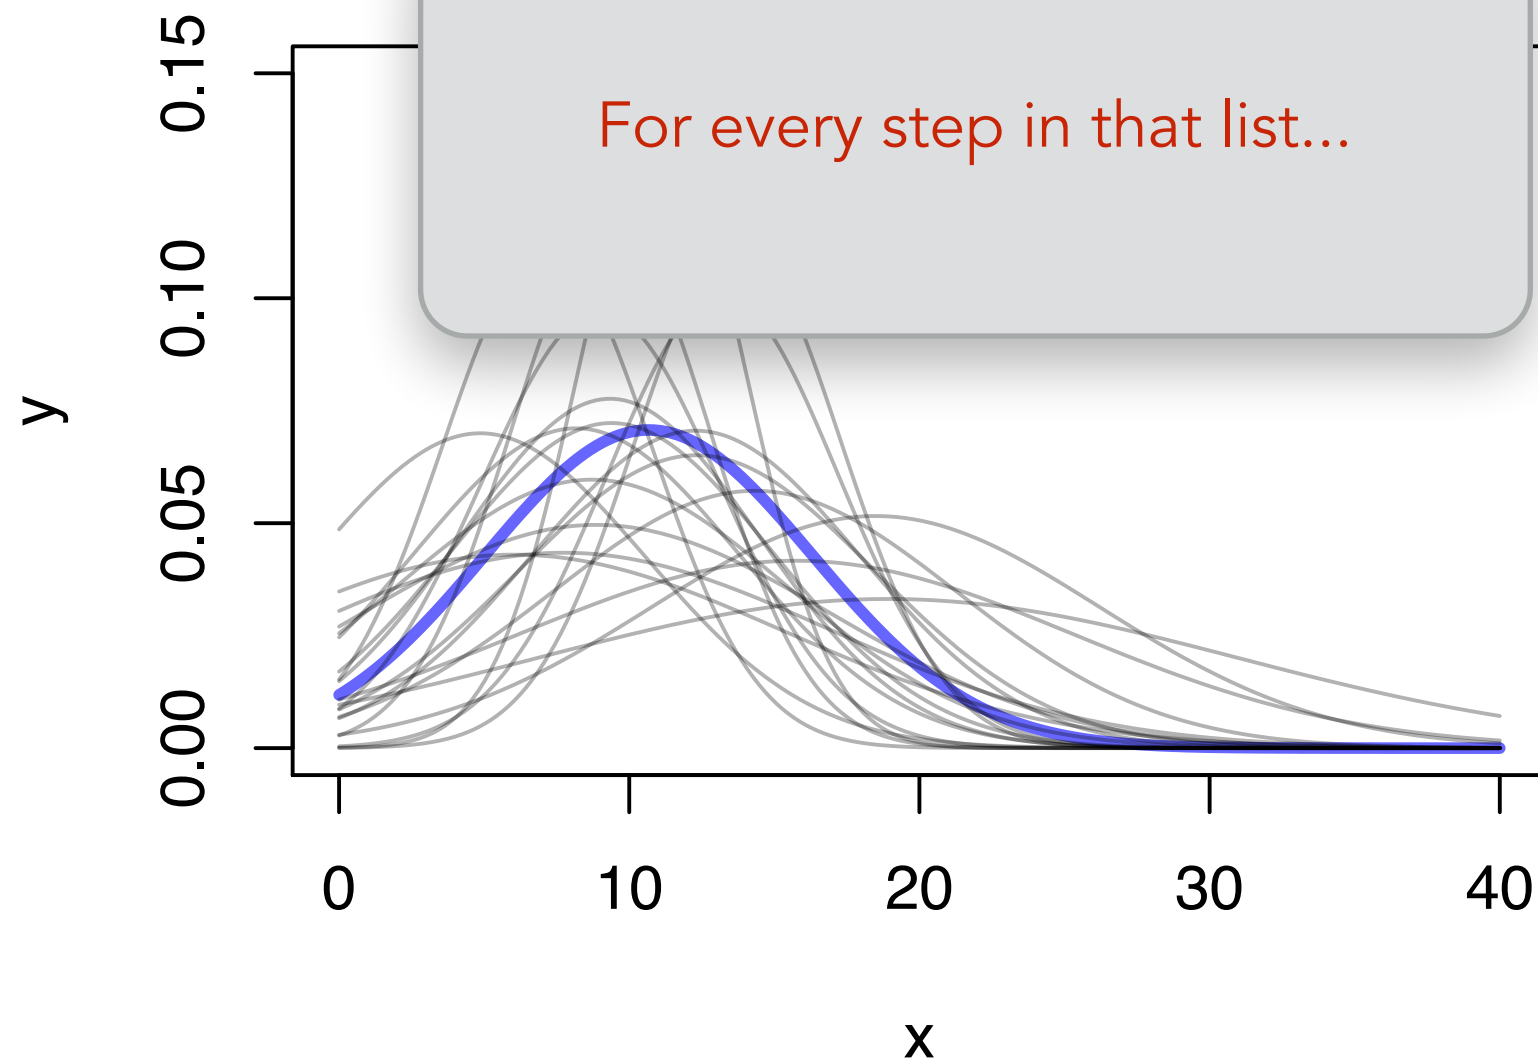
- What about the variation!?

```
samps1 = seq(from = 1, to = length(shopMean), length = 20)
samps2 = round(samps1)
for (i in samps2) {
  y = dnorm(x, mean = shopMean[i], sd = shopMeanSD[i])
  lines(x, y, col = rgb(0, 0, 0, 0.3))
}
```



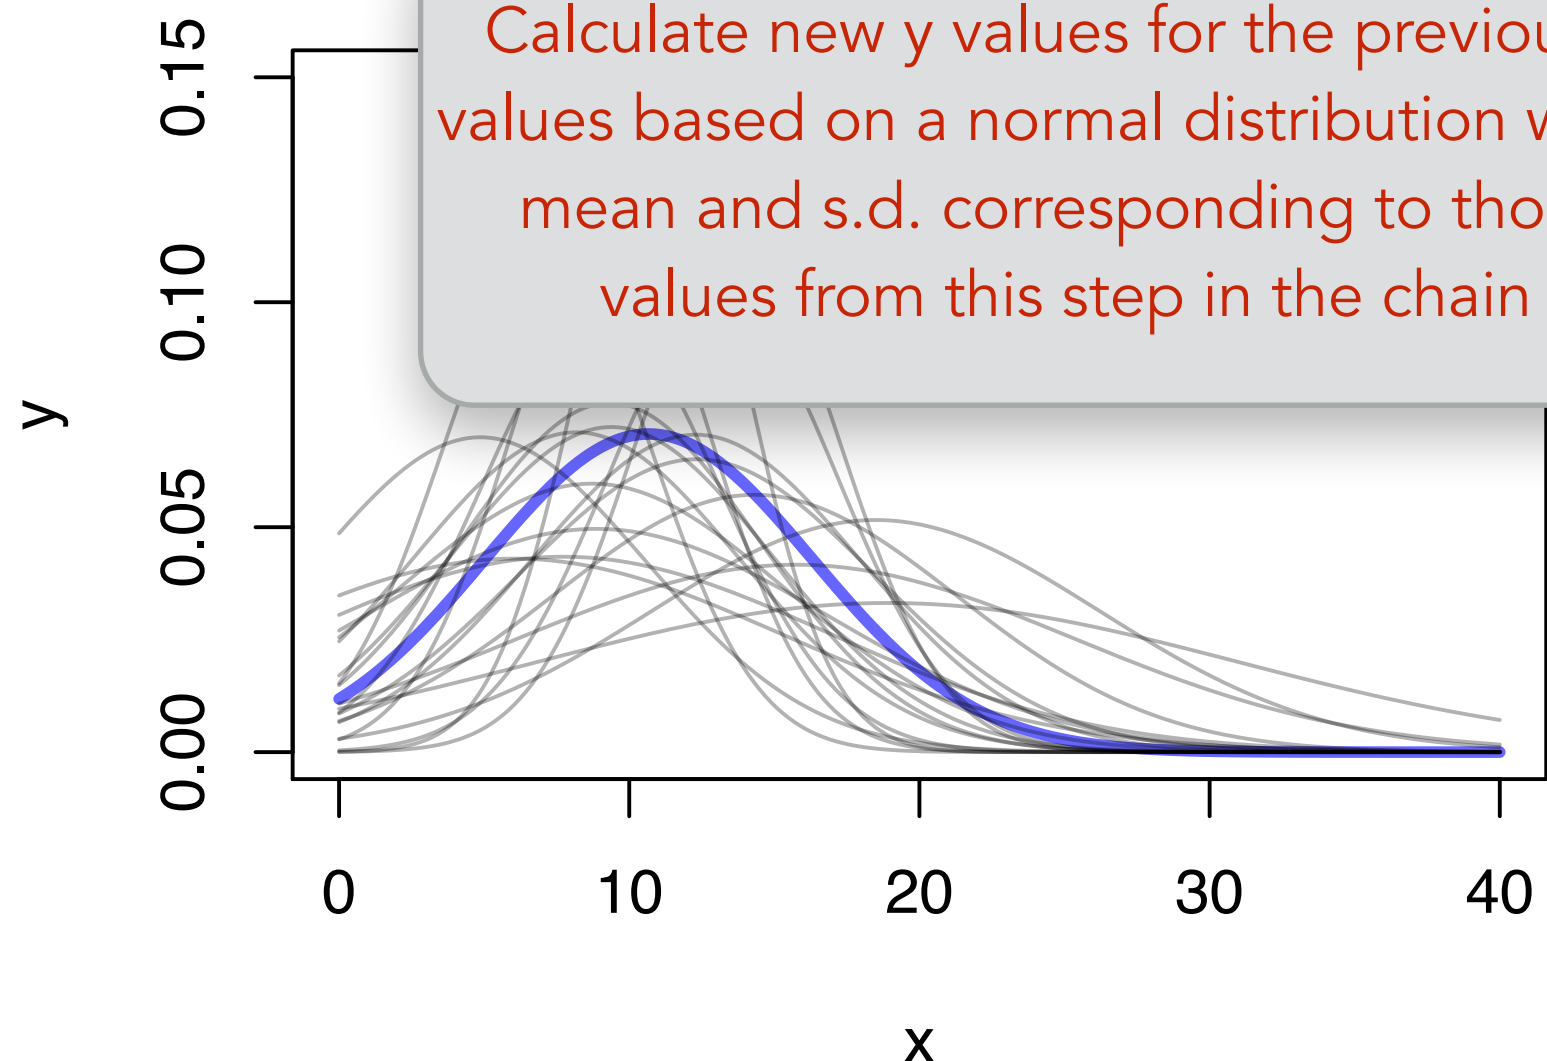
- What about the variation!?

```
samps1 = seq(from = 1, to = length(shopMean), length = 20)
samps2 = round(samps1)
for (i in samps2) {
  y = dnorm(x, mean = shopMean[i], sd = shopMeanSD[i])
  lines(x, y, col = rgb(0, 0, 0, 0.3))
}
```




- What about the variation!?


```
samps1 = seq(from = 1, to = length(shopMean), length = 20)
samps2 = round(samps1)
for (i in samps2) {
  y = dnorm(x, mean = shopMean[i], sd = shopMeanSD[i])
  lines(x, y, col = rgb(0, 0, 0, 0.3))
}
```



- What about the variation!?

```
samps1 = seq(from = 1, to = length(shopMean), length = 20)
samps2 = round(samps1)
for (i in samps2) {
  y = dnorm(x, mean = shopMean[i], sd = shopMeanSD[i])
  lines(x, y, col = rgb(0, 0, 0, 0.3))
}
```





Posterior Predictive Check

- Calculate mean, low, and high expected wait times for each visit


```
#--- Mean expected value for each visit to coffee shop ---#
ypredMean = apply(ypred, 2, mean)

#--- Upper and lower expected 95% HDI for each visit ---#
ypredLow  = apply(ypred, 2, quantile, probs = 0.025)
ypredHigh = apply(ypred, 2, quantile, probs = 0.975)
```

Posterior Predictive Check

- Plot original values


```
par(mfrow = c(1, 1))
plot(coffee$trip, coffee$time, ylim = c(0, 27), xlab = "Trip", ylab = "Wait Time
  (min)", pch = 16, col = rgb(0, 0, 0, 0.6))
abline(v = 30.5)
abline(v = 35.5)
segments(x0 = 0, y0 = 15, x1 = 30.5, y1 = 15, lwd = 2, lty = 2, col = "red")
segments(x0 = 30.5, y0 = 5, x1 = 35.5, y1 = 5, lwd = 2, lty = 2, col = "red")
segments(x0 = 35.5, y0 = 10, x1 = 52, y1 = 10, lwd = 2, lty = 2, col = "red")
```



Posterior Predictive Check

- Add predicted mean values


```
par(new = TRUE)
plot(coffee$trip, ypredMean, ylim = c(0, 27), axes = FALSE, xlab = "", ylab = "",
     main = "", pch = 16, col = rgb(0, 0, 1, 0.5))
```



Posterior Predictive Check

- Add predicted mean values


```
par(new = TRUE)  
plot(coffee$trip, ypredMean, ylim = c(0, 27), axes = FALSE, xlab = "", ylab = "",  
      main = "", pch = 16, col = rgb(0, 0, 1, 0.5))
```



Posterior Predictive Check

- Add predicted mean values


```
par(new = TRUE)  
plot(coffee$trip, ypredMean, ylim = c(0, 27), axes = FALSE, xlab = "", ylab = "",  
      main = "", pch = 16, col = rgb(0, 0, 1, 0.5))
```



Posterior Predictive Check

- Add predicted mean values


```
par(new = TRUE)
plot(coffee$trip, ypredMean, ylim = c(0, 27), axes = FALSE, xlab = "", ylab = "",
     main = "", pch = 16, col = rgb(0, 0, 1, 0.5))
```



Posterior Predictive Check

- Add predicted mean values


```
par(new = TRUE)
plot(coffee$trip, ypredMean, ylim = c(0, 27), axes = FALSE, xlab = "", ylab = "",
     main = "", pch = 16, col = rgb(0, 0, 1, 0.5))
```



Posterior Predictive Check

- Add HDI bars


```
for (i in 1:N) {  
  segments(x0 = coffee$trip[i], y0 = ypredLow[i], x1 = coffee$trip[i], y1 =  
    ypredHigh[i], col = rgb(0, 0, 1, 0.5))  
}
```



Posterior Predictive Check

- Add HDI bars


```
for (i in 1:N) {  
  segments(x0 = coffee$trip[i], y0 = ypredLow[i], x1 = coffee$trip[i], y1 =  
    ypredHigh[i], col = rgb(0, 0, 1, 0.5))  
}
```



Posterior Predictive Check

- Add HDI bars


```
for (i in 1:N) {  
  segments(x0 = coffee$trip[i], y0 = ypredLow[i], x1 = coffee$trip[i], y1 =  
    ypredHigh[i], col = rgb(0, 0, 1, 0.5))  
}
```



Posterior Predictive Check

- Add HDI bars


```
for (i in 1:N) {  
  segments(x0 = coffee$trip[i], y0 = ypredLow[i], x1 = coffee$trip[i], y1 =  
    ypredHigh[i], col = rgb(0, 0, 1, 0.5))  
}
```




Posterior Predictive Check

- Add HDI bars

```
for (i in 1:N) {  
  segments(x0 = coffee$trip[i], y0 = ypredLow[i], x1 = coffee$trip[i], y1 =  
    ypredHigh[i], col = rgb(0, 0, 1, 0.5))  
}
```






Questions?

**Another Application:
Repeated Measurements from
Same Individuals (things)**

Data

- Inter-birth intervals for North Atlantic right whales*
- Multi-locus heterozygosity
- Is there a relationship?



whales x			
	individual	ibi	mlh
1	1	3	0.82
2	1	3	0.82
3	1	4	0.82
4	1	3	0.82
5	2	6	0.54
6	2	6	0.54
7	3	5	0.62
8	4	8	0.40
9	4	9	0.40
10	4	6	0.40
11	4	8	0.40
12	5	3	0.81
13	6	4	0.74
14	7	5	0.59
15	8	5	0.60
16	9	6	0.43
17	10	5	0.60

*Not real data

Photograph used with permission from the New England Aquarium

Data

Some individuals present multiple times...

...others just once

	individual	ibi	mlh
1	1	3	0.82
2	1	3	0.82
3	1	4	0.82
4	1	3	0.82
5	2	6	0.54
6	2	6	0.54
7	3	5	0.62
8	4	8	0.40
9	4	9	0.40
10	4	6	0.40
11	4	8	0.40
12	5	3	0.81
13	6	4	0.74
14	7	5	0.59
15	8	5	0.60
16	9	6	0.43
17	10	5	0.60

Options

1. Ignore that some individuals are repeated, and just analyze the data
 - Incorrect because each data point is not independent

whales x			
	individual	ibi	mlh
1	1	3	0.82
2	1	3	0.82
3	1	4	0.82
4	1	3	0.82
5	2	6	0.54
6	2	6	0.54
7	3	5	0.62
8	4	8	0.40
9	4	9	0.40
10	4	6	0.40
11	4	8	0.40
12	5	3	0.81
13	6	4	0.74
14	7	5	0.59
15	8	5	0.60
16	9	6	0.43
17	10	5	0.60

Options

2. Collapse multiple values to one average for each whale

- Lose a lot of information, and data points

whales x				
Filter				
	individual	ibi	mlh	
1	1	3	0.82	
2	1	3	0.82	
3	1	4	0.82	
4	1	3	0.82	
5	2	6	0.54	
6	2	6	0.54	
7	3	5	0.62	
8	4	8	0.40	
9	4	9	0.40	
10	4	6	0.40	
11	4	8	0.40	
12	5	3	0.81	
13	6	4	0.74	
14	7	5	0.59	
15	8	5	0.60	
16	9	6	0.43	
17	10	5	0.60	

Options

3. Hierarchical model!!!

- Keep all data points, but include “individual” as a categorical parameter in the model
- Make their effects hierarchical
- Accounts for individual effects on the data

whales x				
Filter				
	individual	ibi	mlh	
1	1	3	0.82	
2	1	3	0.82	
3	1	4	0.82	
4	1	3	0.82	
5	2	6	0.54	
6	2	6	0.54	
7	3	5	0.62	
8	4	8	0.40	
9	4	9	0.40	
10	4	6	0.40	
11	4	8	0.40	
12	5	3	0.81	
13	6	4	0.74	
14	7	5	0.59	
15	8	5	0.60	
16	9	6	0.43	
17	10	5	0.60	

Summary

- Should treat nominal data in a hierarchical way by **default**, unless you have a good reason for not doing so
- Hierarchical models are a great way to get the most of your data
 - Don't lose information by averaging across all groups
 - Don't violate assumptions by assuming all independent
- Hierarchical models never hurt (except your brain)
 - Will perform at least as well as a non-hierarchical counterpart
- No limitation to the number of levels
 - Can have groups of groups within groups, etc...
 - As long as you can keep it straight yourself

Questions?